

VI. Distributed Algorithms

Goal: Simulate useful properties of centralistic systems in more or less poorly ordered distributed systems.

c.f.
I-29

[Shingal et al. 1994]: *A distributed system consists of autonomous computers without any shared memory and without a global clock. Computers communicate using message-passing on a communication network with arbitrary delays.* ♦

⇒ **Suitable techniques do not use 'global knowledge'**

Essential: Assumptions w.r.t system model *should be clear!*

► Communication style:

- * point-to-point between single processes
- * broadcast to *all* processes (of a process group)

► Communication guarantees: lost messages, message order, ...

► Which node failures are acceptable for an algorithm?

Reason: *Algorithms do not work without these assumptions!*

Overview: Basic Distributed Algorithms

1. Time and Causality
2. Applications of logical time to message ordering
3. Applications of Time to Distributed Mutual Exclusion and Fairness
4. Consistent global snapshots and checkpointing
5. Determination of 'global' system states: Termination, Deadlocks
6. Distributed Coordination: Leader Election

⇒ **Characteristic techniques for solving distributed problems**

Note: *There are many algorithms/aspects we do not discuss here!*

- * Byzantine Agreement details ⇒ **MSc literature**
- * 2/3 Phase Commit protocols/transactions ⇒ **MSc literature**
- * Algorithms dedicated to unstructured Peer-to-Peer systems
- * Distributed Ledger Algorithms, Bitcoins, Ethereum, ...

VI.1 Time and Causality

Properties of 'real' Time:

linear order: total order relation

Past (linear)/Present/Future (**branching**)

continuous; **dense** \implies **real** numbers as suitable basic model?


Fact: 'Full real time' is not always required

- ▶ traffic lights; backups based on time of day, ... (points in time)
- ▶ Accounting of resource usage (period of time)
- ◀ Decision criteria and priorities for:
 - Resource allocation, e.g., CPU scheduling
 - Election algorithms ('**Eldest**')
- ◀ Version control: text; source-code, e.g., build-organization
- ◀ Synchronization: a_2 *before* a_1 ; not *at the same time*

VI.6.1

\implies **Typically not all properties of real time are really needed.**

Example: Observations and Causality



Wenn ein Zuschauer von der Ferne das Exercieren eines Bataillons verfolgt, so sieht er übereinstimmende Bewegungen desselben plötzlich eintreten, *ehe* er die Commandostimme oder das Hornsignal hört; aber aus seiner Kenntnis der *Causalzusammenhänge* weiß er, daß die Bewegungen die *Wirkung* des gehörten Commandos sind, dieses also jenen *objectiv* vorangehen muß, und er wird sich sofort der Täuschung bewußt, die in der Umkehrung der Zeitfolge in seinen Perceptionen liegt.

Christoph von Sigwart (1830-1904) *Logik* (1889)

Observer: Far away on a hill looking down on military exercising suddenly sees soldiers starting to run and shortly **after** that he also hears the commander who shouts 'RUN'
→ Effect is observable before the cause?

⇒ **Observer time line contradicts rule of cause and effect**

1. $P_{commander}$ sends msg_1 to soldiers (PS)
2. PS reacts, e.g. by running, which is observed by P_{observ} as msg_2
3. P_{observ} hears command msg_1 ⇒ Effect observed before cause?

Rule of Causality: Cause \longrightarrow Effect

- ◁ Observer observes effect before cause
- ◁ **Alibi principle:** Speed of light and the impossibility to be at two distant places at the same time
- ◁ Messages: send is always before the corresponding receive
- ◁ Time paradox in 'backward' time travel
Example: Kill the inventor of the time machine... ?

Note: internal relative order is important, not relation to real time.

\implies **Simple linear orders are sufficient for most CS applications:**

- ▷ Position in FIFO-queue simulates relative arrival time
- ▷ 'Age' of a process simulated by strict monotonous numbering
- ▷ Version control for programs based on ordered Dewey-Notation
- ▷ Mutual exclusion and fairness based on request ordering

\implies **always use the most efficient but sufficient model**

Computer Systems: Real vs. Logical Time

1. 'Real' Time: *Reference to external 'world' and time*

Example: systems that control machines or traffic lights

Physical Clock \approx acceptable deviation from 'real time'

- ▶ internal physical clocks (quartz crystal oscillation)
- ▶ clock alignment within local/global networks
- ▶ external points of reference, e.g. [external time server](#)
⇒ costly in distributed systems if highly accurate

NTP

2. Logical Time: *internal, relative causality-based order*

⇒ reference to real time not needed

Logical Clock \approx internally consistent, no reference to real time

- ▶ integer counter for logical steps (*ticks*)
- ▶ compare and align in the case of message exchange
- ▶ initial time is globally 0 via reset
⇒ cheap and efficient in almost all distributed systems

VI.1.1 Physical Clocks (= real time ?)

- ◀ **Astronomical Time:** $\frac{\text{solar day}}{24 \cdot 60 \cdot 60} = \frac{\text{solar day}}{86400} \approx \text{solar second}$
 Earth/Sun rotation constant, but earth rotation slows down
 \implies solar days/seconds become 'longer' \implies mean value GMT
- ◀ **Atomic time (TAI):** 01.01.1958
 1 solar second = 9.192.631.770 Caesium 133 transitions
- ▶ **Universal Coordinated Time (UTC):** compensates for **Drift**
 currently 3 millisec/day \implies TAI \oplus leap seconds (approx. 0.9 sec)
 \implies lots of problems in IT due to 'repeated second'
- ▶ **Distribution:** accuracy of source \oplus transfer time !
 - Radio-based signals, e.g., DCF77 $\approx \pm 1 \text{ msec} \oplus \pm 10 \text{ msec}$
 - Satellites, e.g., GPS or GNSS (multiple satellites)
 $\approx 10 \text{ nanosec} \rightarrow 1 \text{ microsec}$

Savage
CACM
09/2015
see vc

10^{-9}
 10^{-6}

Caution: 1 nanosec \approx 750 instructions in a 749.070 MIPS processor

Effects of Errors: www.theregister.com/2016/02/03/decommissioned_satellite_software_knocks_out_gps/

AMD/
Rizen 9
3950X
(2019)

Hardware Basis: **Timer Chip**

- Quartz ticks as source for time steps; counter decrement; reset;
- Software clock: **clock ticks** counted via interrupts based on counter
- Deviation ≈ 1 second in approx. $11\frac{1}{2}$ days
clocks diverge in opposite 'direction' \implies approx. 2 seconds max

Algorithms: consistent initialization (*How?*)
keep deviation below threshold by msg exchange

Problems: (in all algorithms)


- ◀ **Transfer time:** Source \longrightarrow Destination
different routes, number of hops, load levels etc.
 \implies measure transfer times and use them in calculations
- ▶ **Suitable methods for adjustment:**
never put clock back; small adjustment instead of abrupt change
 \implies slowdown/accelerate counter for local clock
Expl.: 100 IR/sec \implies 9/11 ms instead of 10 ms as compensation

Approximation of real time – 1

1. Passive Time Server

- (a) Time-Server holds 'external real' time in clock C
- (b) P_i sends Request at T_{snd} and receives C_{server} at T_{rcv}
- (c) $C_{server}^{rcv} \approx C_{server} + \frac{T_{rcv} - T_{snd}}{2}$ (plus service time on server)

Christ-
ian
1989



- ◀ Same time for both messages? (e.g. uni-directional ring)
- ◀ Server and network may have varying loads

Approximation of real time– 2

2. Active Time Server, e.g., Berkeley UNIX

- (a) initially: correct setting of server clock
- (b) **Protocol:** in fixed intervals
 - i. Server sends T_{server} to all (local) network nodes
 - ii. $\forall P_i \in PS$: compute and send $\Delta_i := T_i - T_{server}$ to server by respecting transfer times
 - iii. mean value of divergences is used to correct local clocks

Gusel-
la
et al.
1989

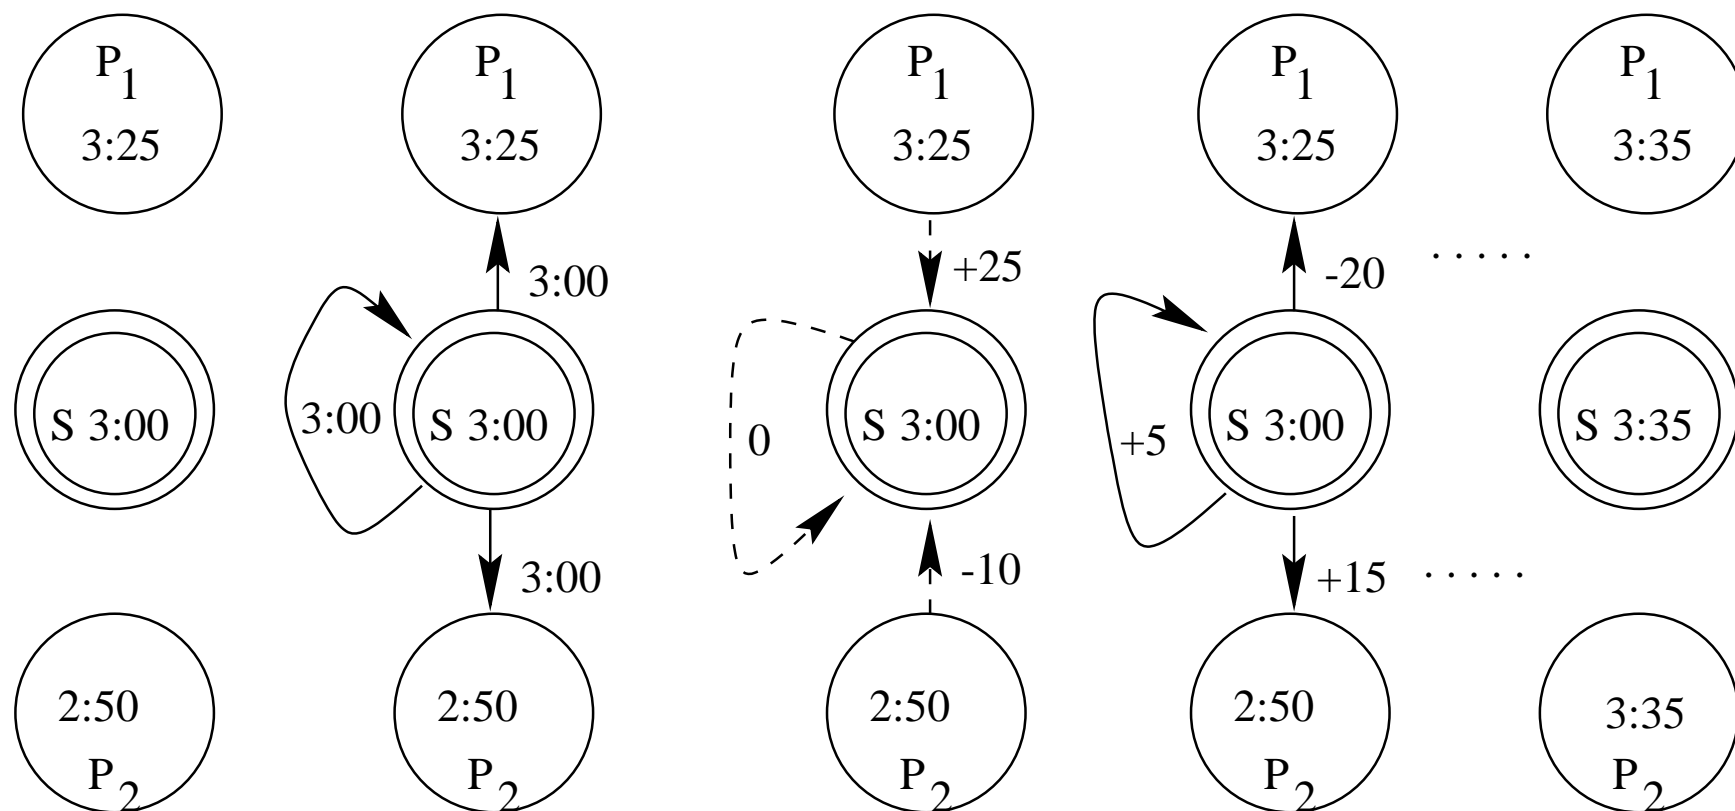
man
timed
c.f.
pg.
VI-11

3. Distributed Adjustment (active)

- Re-Synchronize in fixed (local) intervals
- **broadcast** local time for calibration to all nodes
- compute local mean values for correcting the local clock
requires minimal number of answers; ignores extreme outliers

4. Multiple external clock sources (in different nodes)

Approximation of real time – 3: active Timeserver




Algorithm: compute 'correct' time via arithmetic mean
 send computed deviations
 correction via slowdown/acceleration of local counter

VI.1.2 Logical Clocks - Virtual Time

Cause – Effect relation among

- ▷ Actions or **Events** of the same process
- ▷ Events of different processes due to sending/receiving messages
- ▷ Transitivity



Induced order: (causality relation)

- **local:** $e_{11} \xrightarrow{\sqsubset} e_{12} \xrightarrow{\sqsubset} e_{13} \xrightarrow{\sqsubset} e_{14}$
and $e_{21} \xrightarrow{\sqsubset} e_{22} \xrightarrow{\sqsubset} e_{23} \xrightarrow{\sqsubset} e_{24} \xrightarrow{\sqsubset} e_{25}$
- **Communication:** $e_{12} \xrightarrow{\sqsubset} e_{24}$ and $e_{22} \xrightarrow{\sqsubset} e_{13}$
- **Transitivity**, e.g., $e_{21} \xrightarrow{\sqsubset} e_{14}$
- **concurrent**, e.g., (e_{11}, e_{21}) , (e_{23}, e_{12})

c.f.
III.1

Definition VI.1: (happened-before Relation)

Let a, b, c be events of a set of events E and P_i, P_j processes.

The relation $a \xrightarrow{\sqsubseteq} b$ holds $:\Leftrightarrow$

1. $a, b \in P_i$ and $a \sqsubset b$ in P_i , or (sequential order)
2. $a \approx \text{snd}(\text{msg}, P_j)$ in P_i and $b \approx \text{rcv}(\text{msg}, P_i)$ in P_j , or
3. $a \xrightarrow{\sqsubseteq} c$ and $c \xrightarrow{\sqsubseteq} b \implies a \xrightarrow{\sqsubseteq} b$

The events a and b are

- in a **causality relation** $:\Leftrightarrow (a \xrightarrow{\sqsubseteq} b) \vee (b \xrightarrow{\sqsubseteq} a)$
- **concurrent** $:\Leftrightarrow \neg(a \xrightarrow{\sqsubseteq} b) \wedge \neg(b \xrightarrow{\sqsubseteq} a)$



P and P' **without** communication \implies all events are **concurrent**

- no problem for program logic as there is no interaction
- ◄ no global time among all processes of PS achievable

Lamport's logical clocks

Idea: global time $C : (E, \xrightarrow{\sqsubseteq}) \longrightarrow (\mathbb{N}_0, <)$ for entire PS
respects $\xrightarrow{\sqsubseteq}$ **without** extra **messages**

- $\forall P_i \in PS$ exists a **local counter** C_i initial 0
- $\forall a \in P_i$ exists a local, unique **time stamp** $C_i(a)$
derived from local clock value C_i when executing action a in P_i

Two Rules for global time C ...

C1: $\forall P_i \in PS: a \sqsubseteq_{PS} b$ local in $P_i \implies C_i(a) < C_i(b)$

Time respects local, internal order in each single process

C2: $\forall (a_{snd}, b_{rcv})$ where $a_{snd} \approx \text{snd}(\text{msg}, P_j)$ in P_i and
 $b_{rcv} \approx \text{rcv}(\text{msg}, P_i)$ in P_j
 $\implies C_i(a) < C_j(b)$

Time respects causality between corresponding send/receive

... **ensures:** $a \xrightarrow{\sqsubseteq} b \implies C(a) < C(b)$

Lamport Clocks – Implementation of C1 and C2

IR1: Increment local clock C_i **before** each new action

$\forall P_i \in PS \ \forall a \in P_i$ assign $C_i := C_i + d$ where $(d > 0)$

before executing action a

i.e. $a \sqsubset_{PS} b \implies C_i(b) = C_i(a) + d \implies C_i(a) < C_i(b)$

IR2: Propagate local time information with each message

► Let $a_{snd} \approx \text{snd}(\text{msg}, P_j)$ in P_i

1. C_i is incremented locally to $C_i(a)$ in P_i

2. a_{snd} is **extended** by $\text{snd}(\text{msg}, t_{msg}, P_j)$ where $t_{msg} = C_i(a)$

► Let $b_{rcv} \approx \text{rcv}(\text{msg}, t_{msg}, P_i)$ in P_j

1. C_j is incremented locally to C_{temp} in P_j


2. $C_j := \text{MAX}(C_{temp}, t_{msg} + d)$ where $d > 0$ (transfer time)

i.e. $C_i(a_{snd}) = t_{msg} < t_{msg} + d \leq C_j(b_{rcv})$

Lamport Clocks – Example

$$\triangleright e_{12} \xrightarrow{\sqsubseteq} e_{23}: \text{Max}(\underbrace{2+1}_{rcv}, \underbrace{2+1}_{t_{msg}+d}) = 3$$


$$\triangleright e_{16} \xrightarrow{\sqsubseteq} e_{25}: \text{Max}(\underbrace{4+1}_{rcv}, \underbrace{6+1}_{t_{msg}+d}) = 7$$




Observations:

- ◀ result is a partial order, e.g., (e_{14}, e_{23}) are concurrent
- ◀ $P_i \neq P_j$ use same clock value for different events, e.g., e_{17}/e_{25}
- ◀ concurrent events may have different clock values
Example: (e_{23}, e_{16}) concurrent, but $C(e_{23}) = 3 < C(e_{16}) = 6$

Expl.: 'History' of selected Events in the same PS



Two 'chains' of causal dependencies w.r.t. 'happened-before'



How to define a total order among events?

- ◀ Mapping $C: (E, \xrightarrow{\sqsubset})$ onto $(\mathbb{N}_0, <)$ is not *injective*
- ▶ Based on a unique numbering scheme for all processes $|PS|$:
 $a \xrightarrow{total} b : \iff C_i(a) < C_j(b) \text{ or } (C_i(a) = C_j(b)) \wedge i < j$
- ▶ new mapping from (E, \xrightarrow{total}) onto $(\mathbb{N}_0 \times \mathbb{N}_0, <)$ is injective
- ▶ lexicographical order based on Lamport time and process indices
 \implies **mapping** (E, \xrightarrow{total}) **onto** $(\mathbb{N}_0, <)$ is injective!

Problem: C does not uniquely denote causality!

$$C(a) < C(b) \implies \neg(b \xrightarrow{\sqsubset} a) \quad (\text{not } \iff)$$

- ◀ Increment may be caused locally or by send/receive?
- ◀ Most recent information only w.r.t. sender and receiver

Loss of structural information: $\xrightarrow{\sqsubset} \mapsto <$ and $\xrightarrow{\sqsupset} \mapsto >$
but: $\parallel \mapsto \{<, =, >\}$

Extending Lamport Time to Vector Time

Mat-
tern
1987
Fidge
1988

Idea: Propagate more detailed and also indirect information

- ▶ $|PS| = n \implies \forall P_i \in PS \ \vec{C_i} = \langle C_i[1], \dots, C_i[n] \rangle$
- ▶ $\forall P_i \in PS$ is $\vec{C_i}$ initialized by $\vec{0}$
- ▶ $(i \neq j) \implies \vec{C_i}$ and $\vec{C_j}$ almost always **different**
 - $C_i[i] \approx$ local time C_i in P_i
 - $C_i[j]$ where $(i \neq j) \approx$ **most recent** information in P_i about the local clock value of P_j
 - * a is the most recent action in P_j where $a \xrightarrow{*} b$ holds and
 - * b is the most recent action in P_i
- $\implies C_i[j] = C_j(a) + \Delta$ where $\Delta > 0$
- ▶ **Update:** Messages in PS are extended by so-called **time vectors**
 - \implies only moderate additional overhead w.r.t. Lamport time
 - Information dissemination much faster
 - Processes **without** interaction are not ordered (as before)

Implementing C1 and C2 using Vectors:

c.f.
pg.
VI-14

IR1: Increment the local clock C_i before each new action

$$\forall P_i \in PS \ \forall a \in P_i \text{ assign } C_i[i] := C_i[i] + d \text{ where } (d > 0)$$

$$a \sqsubset_{PS} b \implies C_i[i](b) = C_i[i](a) + d \implies C_i[i](a) < C_i[i](b)$$

IR2: Propagate time vector with each message

- Let $a_{snd} \approx \text{snd}(\text{msg}, P_j)$ in P_i
 1. $C_i[i]$ is incremented locally to $C_i[i](a)$ in P_i
 2. a_{snd} is extended by $\text{snd}(\text{msg}, \overrightarrow{t_{msg}}, P_j)$ where $\overrightarrow{t_{msg}} = \overrightarrow{C_i}$
- Let $b_{rcv} \approx \text{rcv}(\text{msg}, \overrightarrow{t_{msg}}, P_i)$ in P_j
 1. $C_j[j]$ is incremented locally in P_j
 2. $\forall k \in [1 : n] \ C_j[k] := \mathbf{MAX}(C_j[k], \overrightarrow{t_{msg}}[k])$


i.e. corresponding elements of vector $\overrightarrow{C_i}(a_{snd}) \leq \overrightarrow{C_j}(b_{rcv})$

Predicate: $\forall i \in [1 : n] \ \forall j \in [1 : n]$ holds $C_i[i] \geq C_j[i]$

local time in P_i always more recent than its approximation in P_j .

Example: Vector Clocks – 1


- $e_{16} \xrightarrow{\sqsubset} e_{25}$ and $C(e_{16}) = \langle 6, 2 \rangle < \langle 6, 5 \rangle = C(e_{25})$
- (e_{23}, e_{16}) not ordered and
 $C(e_{23}) = \langle 2, 3 \rangle$ **concurrent** $\langle 6, 2 \rangle = C(e_{16})$




Comparison of vector clocks: (for actions a and b)

1. $t^a = t^b \iff \forall k \in [1 : n] \ t^a[k] = t^b[k]$
2. $t^a \leq t^b \iff \forall k \in [1 : n] \ t^a[k] \leq t^b[k]$
3. $t^a < t^b \iff (t^a \leq t^b) \wedge \neg(t^a = t^b)$
4. t^a **concurrent** $t^b \iff \neg(t^a < t^b) \wedge \neg(t^b < t^a)$

Example: Vector Clocks – 2



Example: Comparison of Lamport and Vector Time



Vector time: More information than Lamport

Advantage of vector clocks: $a \xrightarrow{\sqsubseteq} b \iff t^a < t^b$

\implies **Causality can be derived from **time-stamp** alone!**

Reason:

- $t^a < t^b \implies \forall k \in [1 : n]$ holds $t^a[k] \leq t^b[k]$ **and**
 $\exists l \in [1 : n]$ where $t^a[l] < t^b[l]$
- **<-Entries:** local increment distinguishable from message
 \implies **step-wise backtracking** of message transfers
 terminates (**finite** number of processes **and** a-cyclic chain)

Important: 'Logical time' definition depends on algorithm

- *Only events important to an algorithm are time-stamped.*
 e.g., R/W in mutual exclusion, msg snd/rcv in msg ordering, ...
- Size of counters and message overhead is reduced.

Applications of logical time models

c.f.
VI.2/3

Origin: Lamport introduces 'logical time' for fairness in distributed mutual exclusion algorithms. (in: CACM 21(7), 1978)


1. **Message ordering** for point-to-point and broadcast messaging VI.2
2. **Distributed mutual exclusion** VI.3
 - ◁ Request ordering for preventing deadlocks
 - ◁ Request ordering for ensuring fairness
3. Optimistic **Concurrency Control** in distributed database systems
 - conflicts based on mutual dependencies
 - **optimistic** \approx conflicts are assumed to be rare
 - in case of conflict: choose 'victim' process and reset process
 - transactions are **time-stamped** for victim selection
4. ...

End
of
VI.1

VI.2 Message Ordering

Problem: messages do not arrive in the order they were sent

- ◀ Effect: race conditions as in write-write conflicts
Example: data base **updates**; the last update 'wins'
- ◀ Program logic based on causality may fail




Recall: Isolated ordering between exactly two processes is simple in the context of lossless message transfer.

III-65

Ordering of Broadcast-Messages

- **FIFO broadcast:** All **Broadcast-Msg** of a **single** sender P_{snd} arrive for all receivers P_{rcv} in the order as sent.

Example: Problems with a broadcast without FIFO semantics



- **Causal broadcast:** P_{rcv} accepts a **broadcast msg** from P_{snd} **only** after all messages accepted in P_{snd} before sending this **broadcast msg** are also received and accepted in P_{rcv} .

Note: Causal broadcast \implies FIFO broadcast

Causal Broadcast: Birman/Schiper/Stephenson

1991

Preconditions

1. **Broadcast** message transfer is lossless
2. Each $P_i \in PS$ uses a vector for counting message send events

Idea: Time stamps VT for all **broadcast** send actions of all processes

1. Increment $VT_i[i]$ and send updated vector as part of message t_{msg}
2. Use MAX function to update local vector VT_i based on t_{msg}

Test in P_{rcv} : *Compare local VT_i to time stamp vector t_{msg}*

- If **all** local VT_i entries are equal or higher than those in t_{msg}
 - \implies all messages known in P_{snd} have also been accepted in P_{rcv}
 - \implies **accept** incoming message at once
- ◀ If VT_i **misses** messages \implies **buffer** arriving message in P_{rcv} until local vector VT_i is up to date w.r.t. t_{msg} , i.e., P_{snd}

Note: logical time 'counts' broadcast events only

Birman/Schiper/Stephenson-Protocol

1. $\forall P_i \in PS$ initialize \overrightarrow{VT}_i by $\overrightarrow{0}$
 2. **Before** a **broadcast** in $P_i \approx$ increment $VT_i[i]$ locally
 $\implies VT_i[i] = |\text{Messages from } P_i|$
 and send \overrightarrow{VT}_i as part ($\overrightarrow{t_{msg}}$) of the original message
 3. **All** P_j where $(i \neq j)$ receive message including $\overrightarrow{t_{msg}}$
 P_j **buffers** message **until**:
 - (a) $VT_j[i] = t_{msg}[i] - 1 \implies$ message is most recent one from P_i
and
 - (b) $\forall k \in [1 : n] \setminus \{i\}$ holds: $VT_j[k] \geq t_{msg}[k]$
 $\implies P_j$ knows at least all messages known by sender P_i at the time of sending.
 4. **Accept**: Increment $VT_j[i]$ and execute **Buffer-TEST** (3.a/b)
- Note:** logical time \approx Number of messages sent
 acyclic **time stamps** \implies no **Deadlocks** possible in (3.b)

- $\langle 0,0,0 \rangle[1] = \langle 1,0,0 \rangle[1] - 1$
- Other entries in P2 \geq Msg items

- $\langle 0,0,0 \rangle[2] = \langle 1,1,0 \rangle[2] - 1$
- BUT: $\langle 0,0,0 \rangle[1] < \langle 1,1,0 \rangle[1]$
i.e. Msg from P1 not known in P3

- $\langle 0,0,0 \rangle[1] = \langle 1,0,0 \rangle[1] - 1$
- Other entries in $P3 \geq \text{Msg items}$

- $\langle 1, 0, 0 \rangle[2] = \langle 1, 1, 0 \rangle[2] - 1$
- and $\langle 1, 0, 0 \rangle[1] \geq \langle 1, 1, 0 \rangle[1]$

afterwards: $\langle 1, 1, 0 \rangle$

DSG © 2018 . . . 2024 by Guido Wirtz · Distributed Systems Group · WIAI · Otto-Friedrich-Universität Bamberg · Germany VI - 30

Point-to-Point Msg Ordering: Schiper/Egli/Sandoz 1989

Preconditions:

1. **Point-to-Point** message transfer is lossless
2. Each $P_i \in PS$ maintains vector time for all events

Idea: Send local state w.r.t. sending **to all other** processes
 \implies information 'simulates' **broadcast** effect for receiving processes.

Data structure: $PS = \{P_1 \dots, P_n\}$

- **local vector time** \vec{V}_i for all processes $P_i \in PS$
- **Message-List:** Each $P_i \in PS$ uses ML_i to store pairs (P_j, \vec{v}_m)
 holding knowledge about messages sent **to** P_j !!
 $\vec{v}_m \approx$ 'most recent' known vector time when sending in P_j
- **Messages:** $(P_i, \text{msg}, \vec{v}_{msg}, ML_i, P_j)$
 $(P_{snd}, \text{msg}, P_{snd} \text{ time stamp}, ML_i \text{ **without** current msg}, P_{rcv})$

Schipper/Eggli/Sandoz–Protocol – 1

1989

► **Sender** P_i : $\vec{V}_i[i] := \vec{V}_i[i] + 1$;
 $\text{snd}(P_i, \text{msg}, \vec{V}_i, ML_i, P_j)$;
 $ML_i := \text{insert}(ML_i, (P_j, \vec{V}_i))$;
 \implies current message **not** contained in ML_i message list

Order !

◀ **Receiver** P_j : on arrival of $(P_i, \text{msg}, v_{\text{msg}}, \mathbf{ML}, P_j)$ (1)

TEST: if $\nexists (P_j, \vec{v})$ in \mathbf{ML} then ACCEPT; (2) $\in \text{Msg}$

else if $\vec{v} \not\leq \vec{V}_j$ then BUFFER(...); (3)


else ACCEPT; fi; (4)

fi;

(2) (P_j, \vec{v}) in $\mathbf{ML} \approx$ knowledge of P_i about messages to P_j
 pair not found \implies **first message to** P_j (known in P_i)

(3) Condition holds \implies state in P_j is **not** more recent than
 knowledge in P_i w.r.t. msgs to P_j
 when sending \implies Buffer !

Schiper/Eggli/Sandoz-Protocol – 2



ACCEPT: a message $(P_i, \text{msg}, \vec{v}_{msg}, ML, P_j)$ from P_i in P_j

1. Combine $ML \oplus ML_j$ via insert/maximum for entries with $(k \neq j)$
 \implies only the most recent pair (P_k, \vec{v}) remains in ML_j
2. Increment local time \vec{V}_j and compute maximum based on \vec{v}_{msg}
3. Repeat TEST for all messages buffered using new time stamp analogous to steps (2/3) (may result in consecutive ACCEPT steps)

Result: Causal Point-to-Point message ordering

Example – Schiper/Eggli/Sandoz–Protocol

c.f.
pg.
VI-33

1. **snd in** P_1 : increment; $\text{snd}(P_1, \text{msg}, \langle 1, 0, 0 \rangle, \emptyset, P_3)$
afterwards: $ML_1 := [(P_3, \langle 1, 0, 0 \rangle)]$
2. **snd in** P_1 : increment; $\text{snd}(P_1, \text{msg}, \langle 2, 0, 0 \rangle, [(P_3, \langle 1, 0, 0 \rangle)], P_2)$
afterwards: $ML_1 := [(P_2, \langle 2, 0, 0 \rangle), (P_3, \langle 1, 0, 0 \rangle)]$
3. **rcv in** P_2 : the message $(P_1, \text{msg}, \langle 2, 0, 0 \rangle, [(P_3, \langle 1, 0, 0 \rangle)], P_2)$
 $(P_2, ?) \notin [(P_3, \langle 1, 0, 0 \rangle)] \implies \text{ACCEPT}$
 $ML_2 := [(P_3, \langle 1, 0, 0 \rangle)]; \text{increment and MAX} \implies V_2 := \langle 2, 1, 0 \rangle$
4. **snd in** P_2 : increment; $\text{snd}(P_2, \text{msg}, \langle 2, 2, 0 \rangle, [(P_3, \langle 1, 0, 0 \rangle)], P_3)$
afterwards: $ML_2 := [(P_3, \langle 2, 2, 0 \rangle)]$
5. **rcv in** P_3 : the message $(P_2, \text{msg}, \langle 2, 2, 0 \rangle, [(P_3, \langle 1, 0, 0 \rangle)], P_3)$
 $(P_3, \langle 1, 0, 0 \rangle) \in ML$, but $\langle 1, 0, 0 \rangle \not\prec \langle 0, 0, 0 \rangle = V_3 \implies \text{BUFFER}$
6. **rcv in** P_3 : the message $(P_1, \text{msg}, \langle 1, 0, 0 \rangle, \emptyset, P_3)$
 $(P_3, ?) \notin \emptyset \implies \text{ACCEPT}$
Combine \emptyset and $\emptyset \implies ML_3$ remains empty; $V_3 := \langle 1, 0, 1 \rangle$
7. **Buffer check in** P_3 w.r.t. $(P_2, \text{msg}, \langle 2, 2, 0 \rangle, [(P_3, \langle 1, 0, 0 \rangle)], P_3)$
 $(P_3, \langle 1, 0, 0 \rangle) \in ML$, but $\langle 1, 0, 0 \rangle < \langle 1, 0, 1 \rangle = V_3 \implies \text{ACCEPT}$
 ML_3 remains \emptyset ; $V_3 := \langle 2, 2, 2 \rangle$

VI-32
(2)

VI-32
(3)

VI-32
(2)

VI.3 Distributed Mutual Exclusion

Goals:

- Safeness, i.e., correct critical section (*csd*) implementation
- Deadlock freedom, starvation freeness, fairness
- Efficiency: not too much overhead

Def.
III.8
III-25

Model: PS with Interaction using (asynchronous) message-passing

Costs:

- ▶ auxiliary data structures, vector clocks etc.
- ◀ additional processes and/or **messages**

Caution: single point-of-failure \longrightarrow multiple point-of-failure ?

processes do not react, lost control messages ...

\implies **only feasible under strong preconditions**

Logical Time: Fairness and avoiding cyclic waits (Deadlocks)

Approaches to Distributed Mutual Exclusion

- **Centralized approach:** Client/Server model, **but** server may be *bottleneck* \implies hierarchical systems, specialized server
single point of failure \implies replicated, redundant servers c.f. chapter III-66 ff.
 - **Distributed approaches:** Permission to enter *csd* via ... III-71
 - ... **Inquiry among processes:** c.f. pg. VI-41
 - P_i uses message passing to get permission from other processes
 - Permission is considered granted iff majority accepts
 - Inform other processes about own requests and answer requests
 - ... **Exclusive ownership of a control token**
 - Wait or ask for Token using message passing
 - Handover of token after usage or via answering token request
- Variants:** different pre-assumed topologies for message or token exchange with varying overhead etc.

Token-based Algorithms

Basic Idea: Access to *csd* is granted to $P \iff P$ owns **token**

Variants: Organization of token circulation through PS

underlying communication structure of process system PS

Precondition: secure message transfer without loss of control token !!

1. **Simple Algorithm:** PS organized as a logical **ring** $P \mapsto P_{next}$

`rcv(token)` \implies

`IF (request(self)) THEN csd ELSE snd(P_{next} , token); FI;`

Problems: 'unused' token circulates around the ring

node **crashes**; **lost messages**

2. **Suzuki-Kasami Broadcast Algorithm (1985)**

▷ P_i wants *csd* $\implies P_i$ sends **broadcast** with *REQ* for token

▷ Owner of token reacts only on incoming *REQ*

▷ *REQ* messages are prioritized based on logical *REQ* counter

\implies Token handover after *csd* acts fair based on time stamps

Suzuki-Kasami Broadcast Algorithm

Data structures: (REQ -counter as vector time)

- local $RN_i[i]$ counter for most recent request REQ in P_i
- Vector $RN_i[1 : |PS|] \approx$ most recent, known REQ of other P_j
- **Token: Queue Q holding requesting processes and**
Vector $LN[1 : |PS|]$ holding numbers of **granted REQ**

Problems to be solved:

- ◀ **Outdated REQ** should not be answered

Example: All processes receive REQ from P_1 that is granted by P_2 ; afterwards, other processes should not hand over to P_1 again
 \implies **granted** requests LN are part of the token

- ▶ **Decide for next process to hand over:** (Fairness)

Queue Q of all processes with pending REQ is part of the token
and updated after each csd access

Suzuki-Kasami Broadcast Algorithm

► Request in P_i :

```

 $RN_i[i] := RN_i[i] + 1;$                                      /* Prologue */
broadcast( $i, REQ, RN_i[i]$ ); rcv( $Q, LN$ );      /* to all procs including  $P_i$  */
      csd                                           /* blocking rcv: wait for token ... csd */
 $LN[i] := RN_i[i];$                                            /* Start of Epilogue */
FORALL  $j \in [1 : n]$  DO                                     /* Token-Q update */
    IF ( $RN_i[j] == LN[j] + 1$  AND  $P_j \notin Q$ ) THEN  $Q.enqueue(P_j)$ ; FI;
    OD;
IF ( $Q.notempty()$ ) THEN snd( $Q.frontdequeue(), Q, LN$ );    FI;

```

► Reactions in P_i :

```

rcv( $j, REQ, RN_r$ )  $\implies$                                      /* always able to react */
     $RN_i[j] := MAX(RN_i[j], RN_r);$                              /* update REQ counter */
    IF (token() AND (NOT in csd/Epilog) AND ( $RN_i[j] == LN[j] + 1$ ))
    THEN snd( $P_j, Q, LN$ );
    FI;                                                         /* empty Q due to Epilogue ! */

```

Properties of Suzuki-Kasami Broadcast Algorithm

- ▶ **Safeness:** At any time, token is owned by at most one process
entering csd only possible when owning token
- ▶ **Fairness:**
 - $(n + 1)$ -th *REQ* enters Q only if n -th request has succeeded
 - \implies **maximum** number of $(|PS|-1)$ processes in **FIFO-Queue** Q ;
 - P hands token to next process after csd if there is a *REQ* in Q
- ◁ If P_i owns token and there are **no** *REQ* in Q
 - $\implies P_i$ may use csd repeatedly
- ▶ **Blocking:** Token is owned by P_i after a finite time (c.f. Fairness)
but: crashes, lost messages, non-terminating csds
- ▶ **Cost:** $0 \dots |PS|$ messages (*REQs* and token)

Simulation of central knowledge (server):

Current owner of token acts as the server for the csd.

Non-Token Algorithms

Preconditions:

- ▶ global, unique time stamps for all messages
- ▶ message transfer respects message ordering !

Structure: each process $P_i \in PS$ 'knows' about it's process sets

1. **Request set** $R_i \subseteq PS$: **ask** for permission to enter *csd*
2. **Inform set** $I_i \subseteq PS$: **notify** if local state w.r.t. *csd* changes

Idea: minimize sets for each P_i

- R_i : ensure that at most one process is able to enter *csd*
- P_i is part of a sufficient number of inform sets I_j of other processes
 $P_j \implies P_i$ has a solid basis for it's own decisions
- in case of conflict: minimal time stamp \implies highest priority

Variants: Size of R_i and I_i

Number and content of messages required

Lamport Algorithm

CACM
21(7),
1978

Remark: First application of logical time concepts

- ▶ Simulate 'global server' through **all** processes of PS
 - **Request set:** $\forall P_i \in PS$ choose $R_i = PS \setminus \{P_i\}$
 - **Inform set:** $\forall P_i \in PS$ choose $I_i = PS \setminus \{P_i\}$
 \implies **global message exchange in PS**
- ▶ **Data structures** in each P_i : **Priority-Queue** Q_i holding request pairs
 - $\langle t_j, j \rangle$ lexicographically ordered by **time stamps** in Queue
 - process numbers and Lamport clocks $\approx clock()$
 \implies globally ordered time stamps t_i
- ▶ **Message types:** $REQ \approx$ request
 $REPLY \approx$ acknowledgement for REQ
 $REL \approx$ release

c.f.
pg.
VI-18

```


► Request in  $P_i$ :  $t_i \approx \text{clock}()$ ; /* Prologue */
FORALL  $P_j \in R_i$  DO  $\text{snd}(P_j, \langle REQ, t_i, i \rangle)$ ; OD; /* inform */
 $Q_i.\text{enqueue}(\langle t_i, i \rangle)$ ;
WAIT UNTIL /* periodic test */
    FORALL  $P_j \in R_i$   $\text{rcv}(P_j, \langle REPLY, t_j, j \rangle)$ ; (1)
    AND ( $Q_i.\text{front}() == \langle t_i, i \rangle$ ); (2)
csdi;
 $Q_i.\text{dequeue}()$ ; /* Epilogue */
FORALL  $P_j \in I_i$  DO  $\text{snd}(P_j, \langle REL, t_i, i \rangle)$ ; OD;

```



$$\begin{aligned} \text{rcv}(P_j, \langle REQ, t_j, j \rangle) &\implies \text{snd}(P_j, \langle REPLY, t_i, i \rangle); \\ &\quad Q_i.\text{enqueue}(\langle t_j, j \rangle); \\ \text{rcv}(P_j, \langle REL, t_j, j \rangle) &\implies Q_i.\text{remove}(\langle t_j, j \rangle); \end{aligned} \tag{3}$$

Example: Lamport Algorithm - Two Requests


Le-
gend
c.f.
pg.
VI-46



Example: Lamport Algorithm - Permission & Wait

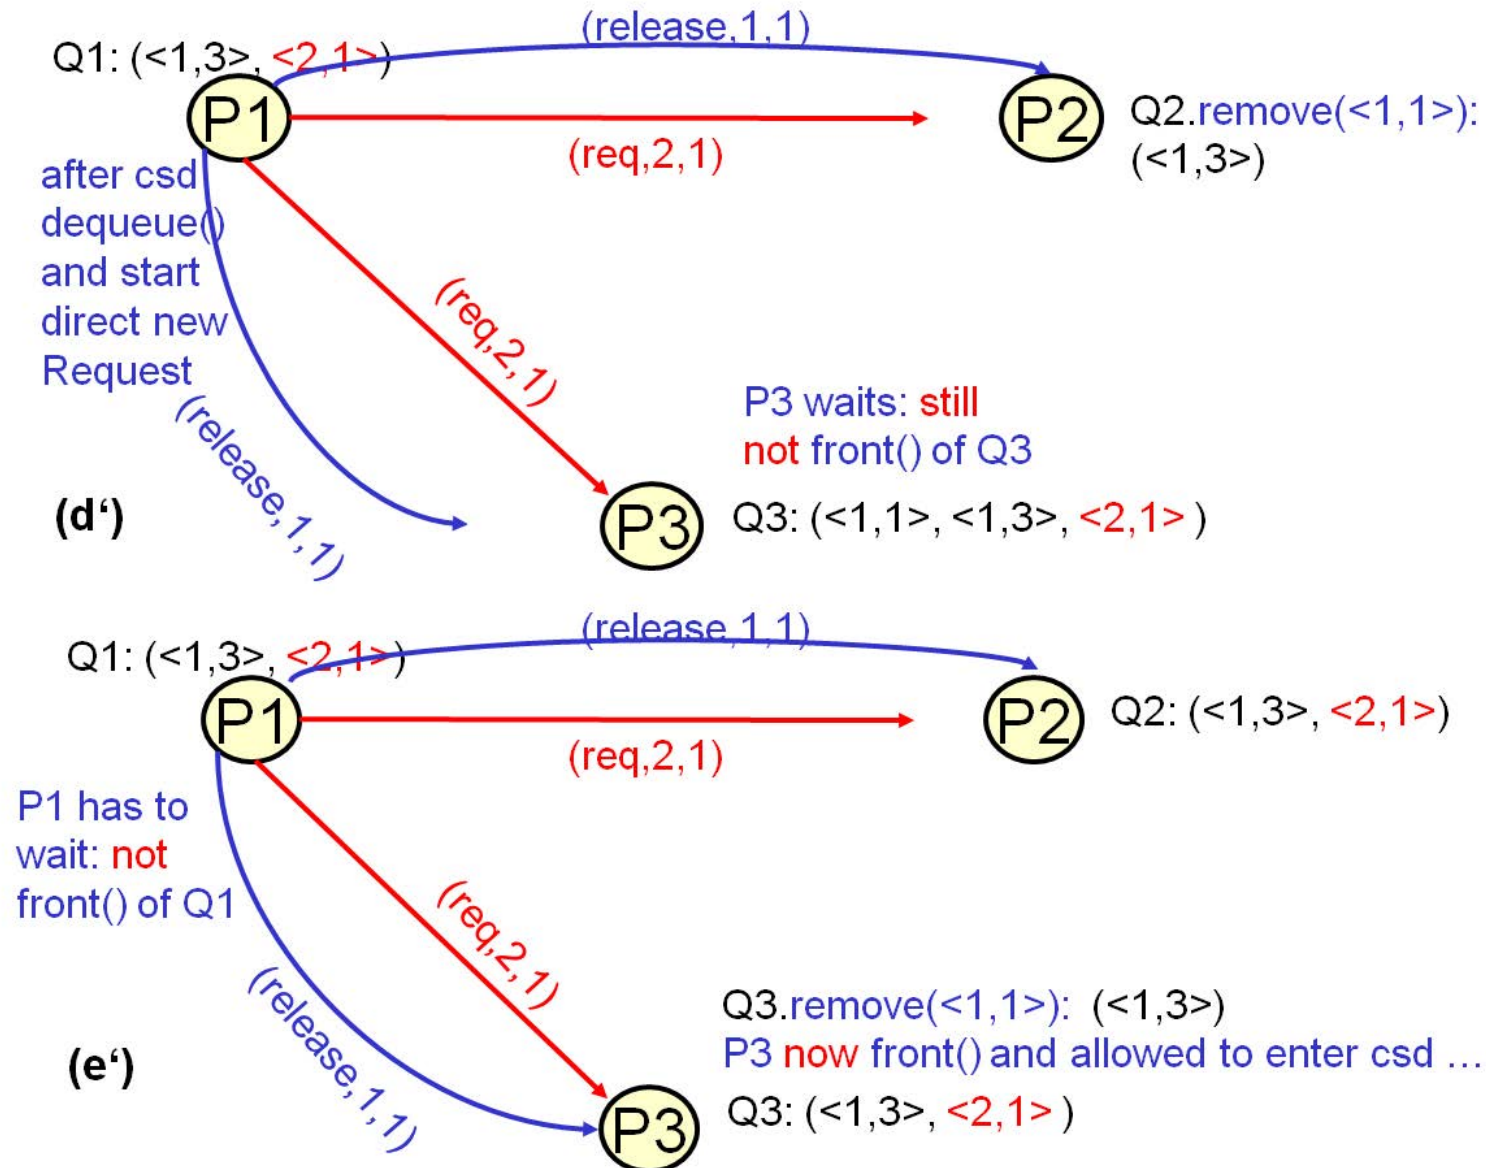


Example: Lamport Algorithm - Handover



Note: Request counter initial 0 in all processes;
 before sending Request: increment by 1
 All queues are empty initially
 Messages carry 3 items of information (<Type>, ReqNr, ProcID)
 where <Type> ::= req | reply | rel

Example: Lamport Algorithm - Unfairness ?



Assessment of Lamport Algorithm

- ▶ Safeness: P_i in $\text{csd}_i \implies$ Wait actions (1) and (2) executed
 - (1) REQ ; wait for $REPLY \implies$ processes know about REQ_i
 - (2) own request is front \implies other requests are more recent
 \implies queues of all other processes wait for REL from P_i
- ▶ No permanent blocking: (3) ensures that all $REPLY$ s are sent
- ▶ No deadlock: globally ordered time stamps prevent from cycles
- ▶ Fairness: conflicting requests are handled fair by REQ order queues
- ◀ **Cost:** $3 * (|PS| - 1)$ message for each csd permission
 \implies **optimization advisable**
- ◀ **Reliability:** *realistic preconditions?*
 - entire algorithm does not work if a single process fails to answer
 - lost $REQ/REPLY$ blocks P_i ; lost REL blocks csd
 \implies Use **time outs** to detect faulty processes

no
lost
msgs

Optimization: Ricart-Agrawala Algorithm

CACM
1981

Basic Idea: combine *REPLY* and *REL* messages

► Request in P_i :

```

FORALL  $P_j \in R_i$  DO  $\text{snd}(P_j, \langle REQ, t_i, i \rangle)$ ; OD;           /* inform */
WAIT FORALL  $P_j \in R_i$   $\text{rcv}(P_j, \langle REPLY, t_j, j \rangle)$ ;
 $\text{csd}_i$ ;
FORALL  $P_j \in Q_i$  DO  $\text{snd}(P_j, \langle REPLY, t_i, i \rangle)$ ;           /*Epilogue*/
 $Q_i.\text{remove}$  OD;

```

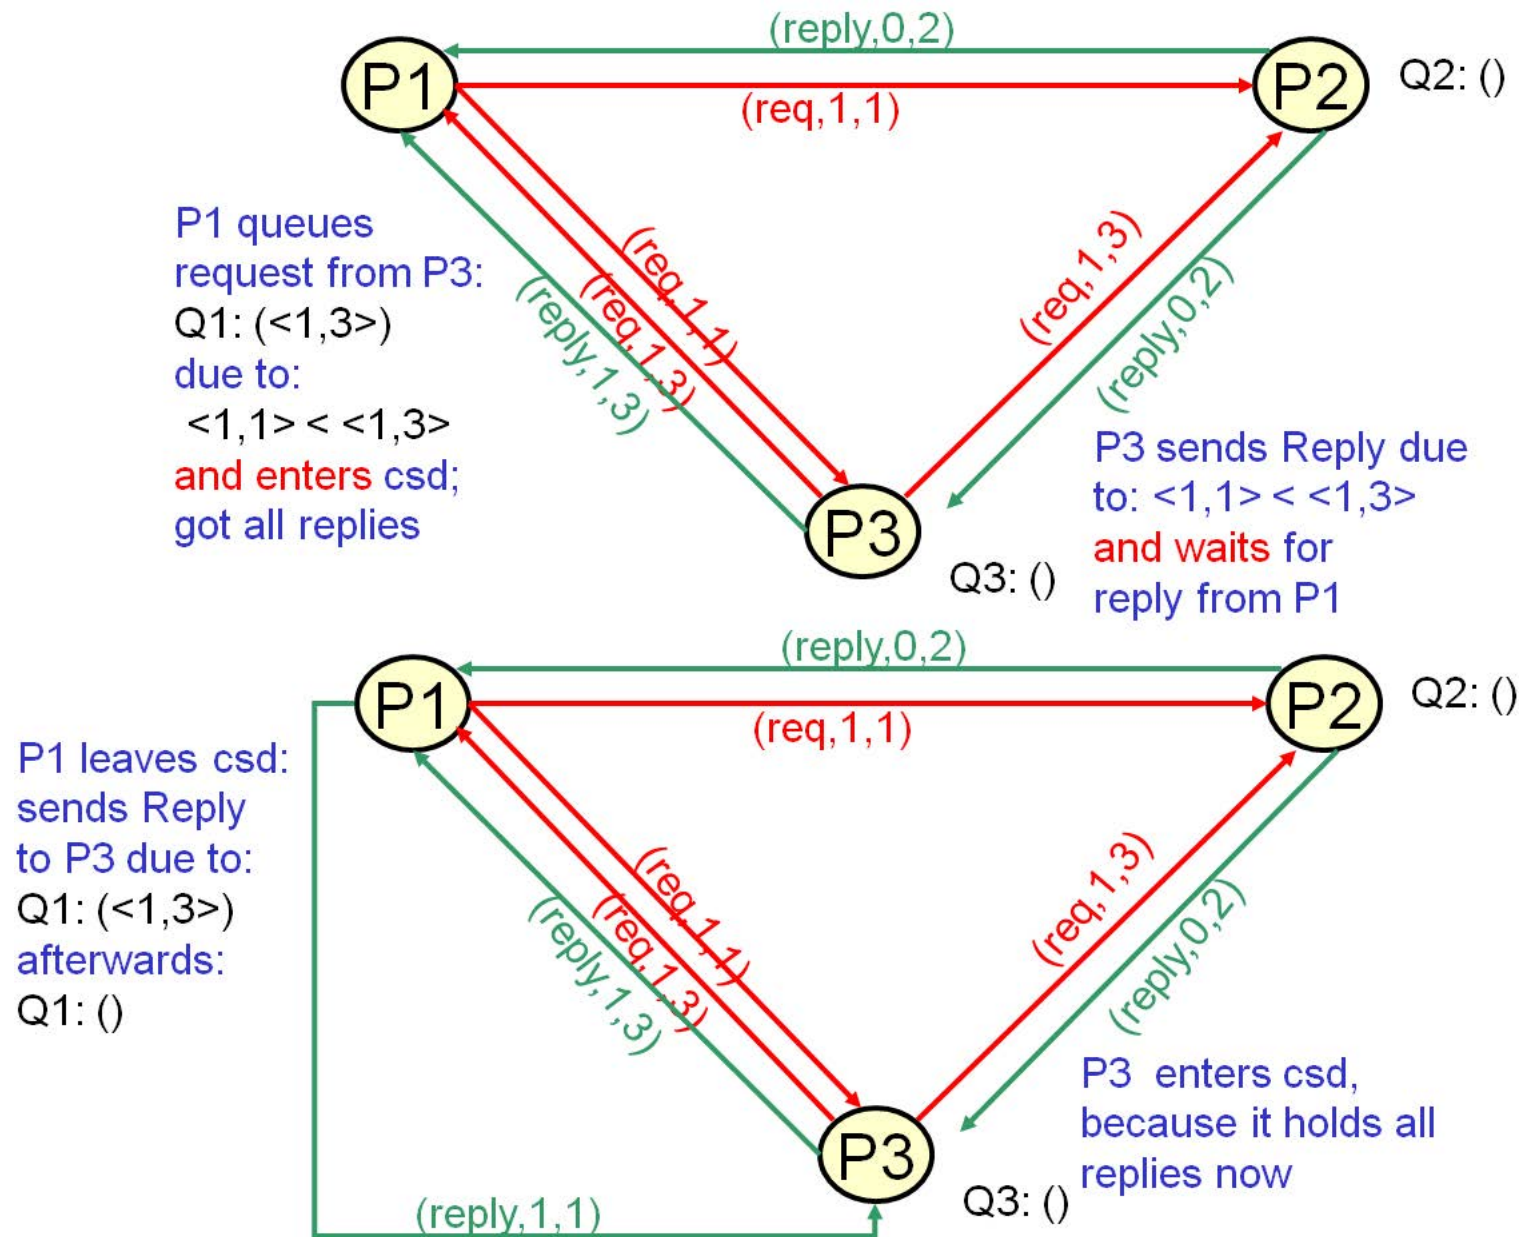
► Reactive behavior in P_i at all times:

```

 $\text{rcv}(P_j, \langle REQ, t_j, j \rangle) \implies$ 
  IF [ ( $P_i$  not in  $\text{csd}$ ) AND ( $t_j < t_i$ ) ]           /* csd not needed */
    THEN  $\text{snd}(P_j, \langle REPLY, t_i, i \rangle)$ ;           /* or 'older' request */
  ELSE  $Q_i.\text{enqueue}(\langle t_j, j \rangle)$ ;

```

Example: Ricart-Agrawala Algorithm



Assessment of Ricart-Agrawala Algorithm

- Process stores requests only if it owns the *csd*
 - ▷ unburdens processes that don't need *csd* at all
 - ◁ not much redundancy for storing requests
- all processes have to react and send *REPLY* messages
- when leaving cs_i all pending reply messages are send

Safeness: P_i waits for **all** *REPLY* messages (as before)
only processes P_j that are not inside the *csd*
or have only a lower priority request will answer

Cost: $2 * (|PS| - 1)$ messages for each *csd* permission

Note: *There are a lot more algorithms and optimizations dealing with distributed mutual exclusion. An overview can be found in P. C. Saxena and J. Rai: A survey of permission-based distributed mutual exclusion algorithms. Computer Standards & Interfaces, 25(2), 2003*

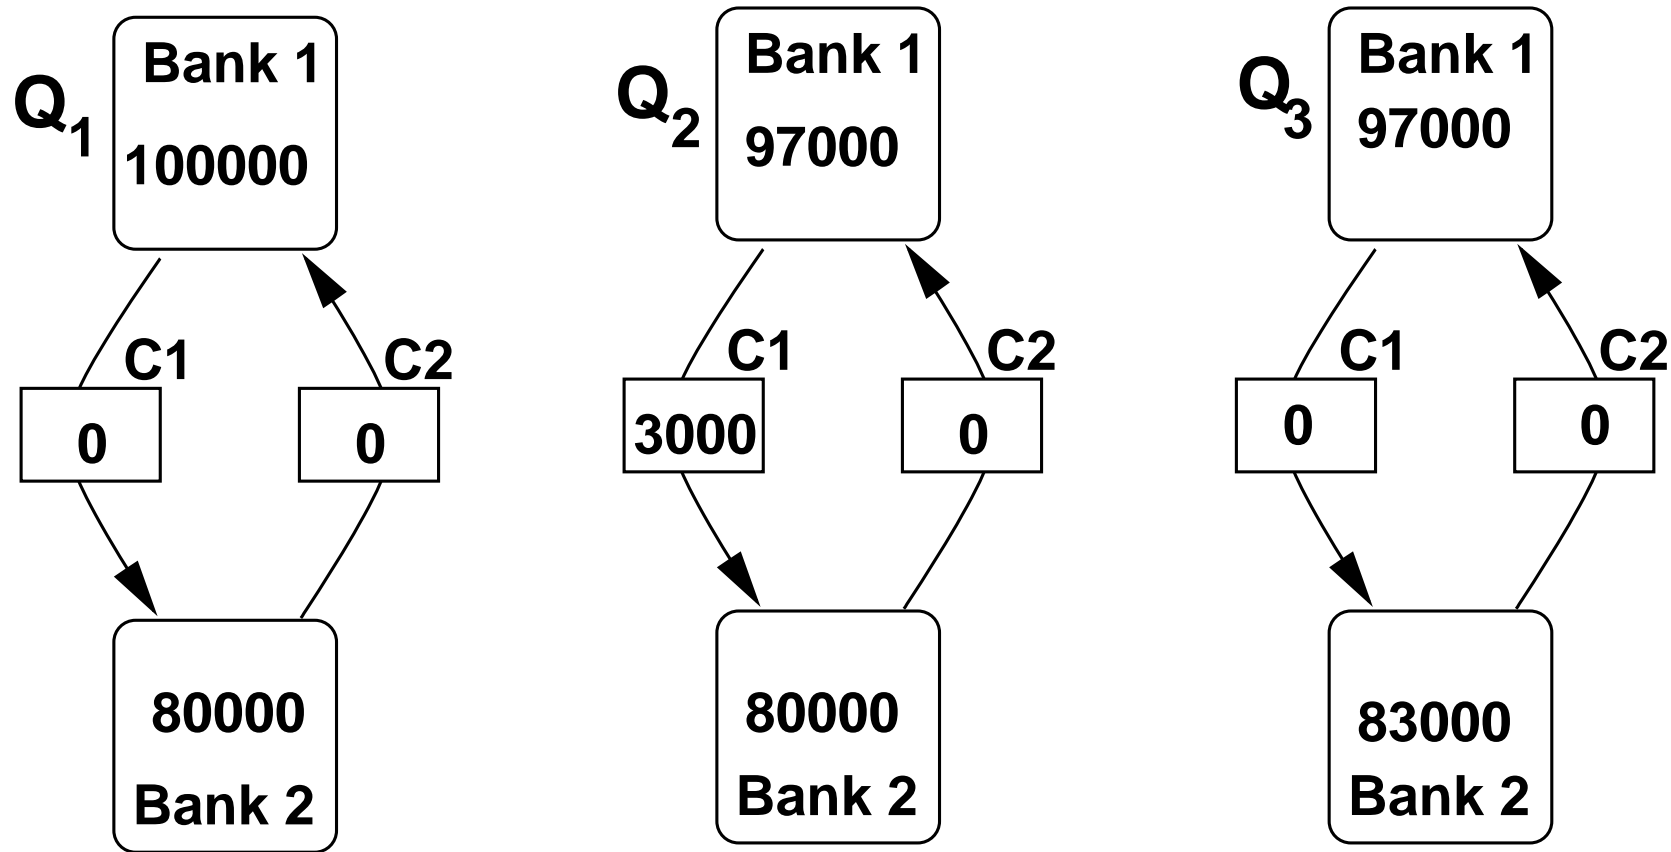
VI.4 Global Snapshots and Consistency

Motivation: **Fault Tolerance**

- ▶ **Why:** many compute nodes and complex networks
 \implies **high probability of (partial) failures**
 - ▶ **How:**
 1. Store **local states** of processes
 Examples: *MEM*, Register, *PC*, Resources
 2. Store the content of **message channels**
 3. **Combine** distributed local states to global **Recovery Points**
 4. Crash \implies **Roll-back** based on most recent recovery point(s)
 - ◀ **Problems:** additional overhead
 - **Trade-Off:** Memory/Communication vs. **Rollback** benefits
 - **Consistency** of stored 'global' states!
- Overhead at runtime \implies Crash easier to handle**

PCB
cf.
II-17

Example: Consistency is an Important Issue



- global state \approx amount of money in (bank1, C1, C2, bank2)
- **inconsistent state recording without coordination, e.g.,**
 1. Bank1 in Q_1 , channels/bank2 in Q_2 : (100000,3000,0,80000) \approx 183000
 2. Channels in Q_1 , bank1/2 in Q_2 : (97000,0,0,80000) \approx 177000

Reasons for Consistency Problems

◀ Processes:

- uncoordinated recording of local states is not sufficient
- coordination based on 'global clock' usually not feasible

Combination: internal state **plus** view on external system

◀ Message channels: How to get content?

- record content of channels (before/after sending) **or**
- wait for channel clearance based on maximum transfer time

Important: (P_1, C_{12}, P_2) observe $|SND_{P_1}| = |C_{12}| + |RCV_{P_2}|$
without FIFO channels \implies explicit msg ordering required

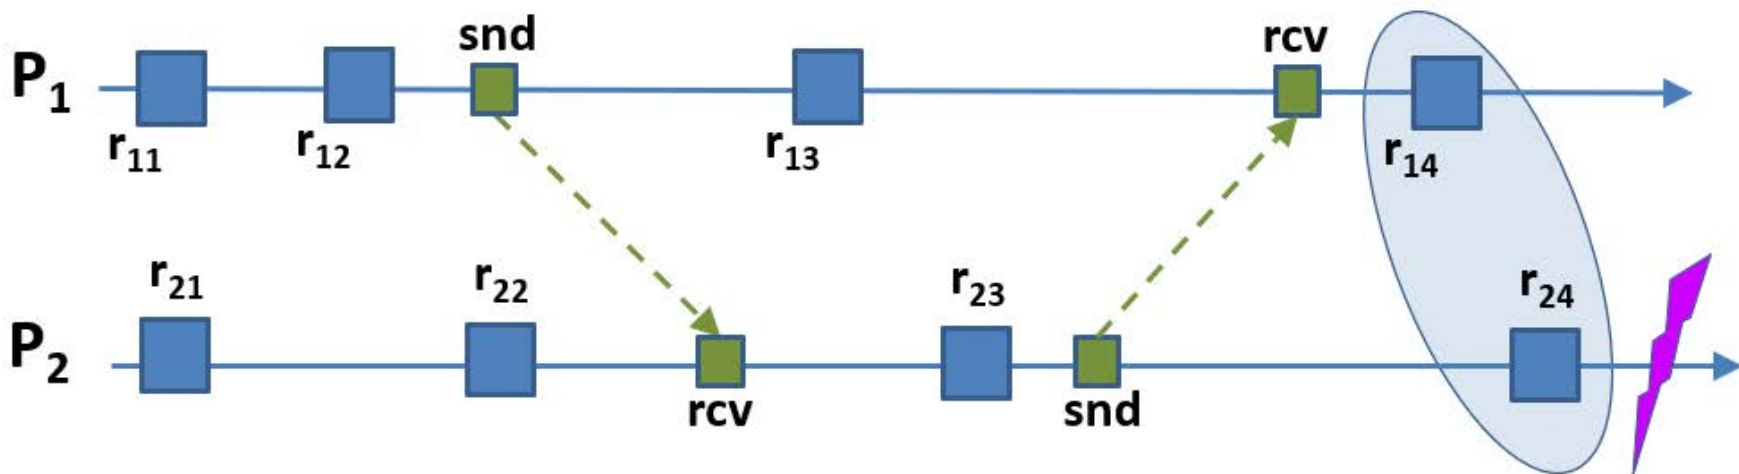
Elementary Problem: $P_i \neq P_j$ record **different** external views

- ◁ wrong assumptions about messages sent
- ◁ wrong assumptions about messages received
- ◁ how to detect messages lost during transfer?

Problems using local Recovery Points only – 1

- Each process P_i records its own local recovery-points r_{ik}
- Recording is **not** coordinated among processes
- Crash $\implies PS$ resets each P_i to its *most recent* recovery-point
- Messages are re-send if sent after recording recovery-point


1. Crash in $P_2 \implies$ **roll-back** to state (r_{14}, r_{24})




\implies Useful Procedure that keeps lot of information? Not always!

c.f.
pg.
VI-57

Problems using local Recovery Points only – 2



2. Crash in P_2 : (r_{13}, r_{22}) ? \implies **lost message** (will block P_2)
 Reset P_1 to r_{12} \implies **roll-back** to state (r_{12}, r_{22})




3. Crash in P_1 : (r_{14}, r_{23}) ? \implies **orphan message** (repeated snd)
 Reset P_1 to r_{13} \implies **roll-back** to state (r_{13}, r_{23})

Uncoordinated Procedure leads to Domino Effect

$\forall (P_i, P_j)$ check whether $\approx |\text{SND}_{ij}| = |\text{RCV}_{ji}|$?

- **Lost Messages** $|\text{SND}_{ij}| > |\text{RCV}_{ji}| \implies$ Reset SND process
 - **Orphan Messages** $|\text{SND}_{ij}| < |\text{RCV}_{ji}| \implies$ Reset RCV process
- until global recovery-point with *consistent information* is reached.



- Crash in $P_2 \implies (r_{15}, r_{24})$ includes Orphan 1 \implies reset P_2
 - (r_{15}, r_{23}) includes Orphan 2 \implies reset P_1
 - (r_{14}, r_{23}) includes lost message 3 \implies reset P_1
 - (r_{13}, r_{23}) includes lost message 4 \implies reset P_2
 - (r_{13}, r_{22}) includes Orphan 5 \implies reset P_1
 - (r_{12}, r_{22}) includes Orphan 6 \implies reset $P_2 \implies (r_{12}, r_{21})$
- Domino-Effect**
(worst-case)

Coordinated Procedure – General Considerations

⇒ **Detect and avoid inconsistencies when recording**

1. Lost Messages: Record contents of **message channels**

- ▶ fault-tolerance in network protocols and OS network interface
e.g., store-and-forward until successful acknowledgement
- ▶ repeating a SND is no big deal!

⇒ **Do not roll-back if** $|SND_{ij}| > |RCV_{ji}|$

2. Orphan Messages: repeated SND may disrupt receiver process avoid this kind of *inconsistency*

⇒ **Do rollback if** $|SND_{ij}| < |RCV_{ji}|$

Coordinated Approach:

- Coordinate local state recordings in order to avoid *domino effect*
- Provide **secure message channels**
 - ▶ System **active**: $\forall (P_i, C_{ij}, P_j)$ holds $|SND_{P_i}| = |C_{ij}| + |RCV_{P_j}|$
 - ▶ System **inactive**: $\forall (P_i, C_{ij}, P_j)$ holds $|SND_{P_i}| = |RCV_{P_j}|$

empty
channels

System Model for Snapshot Algorithm – 1

- ▶ Process system $PS = \{P_1, \dots, P_n\}$
- ▶ **Complete Connectivity:** among each pair (P_i, P_j) of processes there exists a message channel C_{ij}
 1. Messages are **not** lost: middle-ware layer for lossless messaging
 2. Message channels have FIFO property, i.e., all channels (P_i, P_j) support proper message ordering.
- ▶ All processes and channels are *observable* (in theory)

Notation:

- $\forall P_i \in PS$ let LP_i be the local state of P_i
- **Observed Actions:** Send/Receive/Record state
 - * $\text{snd}(m_{ij})$ sending the message m_{ij} in P_i
 - * $\text{rcv}(m_{ij})$ receiving the message m_{ij} in P_j
 - * LP_i recording the local state LP_i in P_i
- $\text{time}_i(\text{event}) \approx$ local *time* in P_i when *event* occurs

System Model for Snapshot Algorithm – 2

Basic Principles:

1. Which snd/rcv events are part of recording a local state:

$$(a) \text{ snd}(m_{ij}) \in LP_i : \Longleftrightarrow \text{time}_i(\text{snd}(m_{ij})) < \text{time}_i(LP_i)$$

$$(b) \text{ rcv}(m_{ij}) \in LP_j : \Longleftrightarrow \text{time}_j(\text{rcv}(m_{ij})) < \text{time}_j(LP_j)$$

\implies local state recording respects **causality** due to sequential execution order in P_i .

2. Compare local states of different processes P_i and P_j ($i \neq j$):

$$(a) \text{ TRANSIT}(LP_i, LP_j) :=$$

$$\{ m_{ij} \mid \text{snd}(m_{ij}) \in LP_i \wedge \text{rcv}(m_{ij}) \notin LP_j \}$$

msgs sent but not received w.r.t. compared states of P_i and P_j

$$(b) \text{ INCONSISTENT}(LP_i, LP_j) :=$$

$$\{ m_{ij} \mid \text{snd}(m_{ij}) \notin LP_i \wedge \text{rcv}(m_{ij}) \in LP_j \}$$

msgs received but not sent w.r.t compared states, i.e. **Orphans**

System Model for Snapshot Algorithm – 3

Characteristics of global states:

1. **Global State** $GS := \{ LP_1, LP_2, \dots, LP_n \}$ such that holds:

$$\forall P_i \in PS \exists \text{ exactly one } LP_i \in GS$$
 isolated recording of one local state for each process in PS
2. GS is **consistent** $:\iff GS$ contains **no Orphans**, i.e.,

$$\forall i \in [1 : n] \forall j \in [1 : n] \text{ holds: } \text{INCONSISTENT}(LP_i, LP_j) = \emptyset$$
3. GS is **strong consistent** $:\iff GS$ consistent **and**

$$\forall i \in [1 : n] \forall j \in [1 : n] \text{ holds: } \text{TRANSIT}(LP_i, LP_j) = \emptyset$$
 combined local states hold all information due to empty channels

Note: If GS is consistent **and** there is an upper bound for all msg transfer times t_{max} :

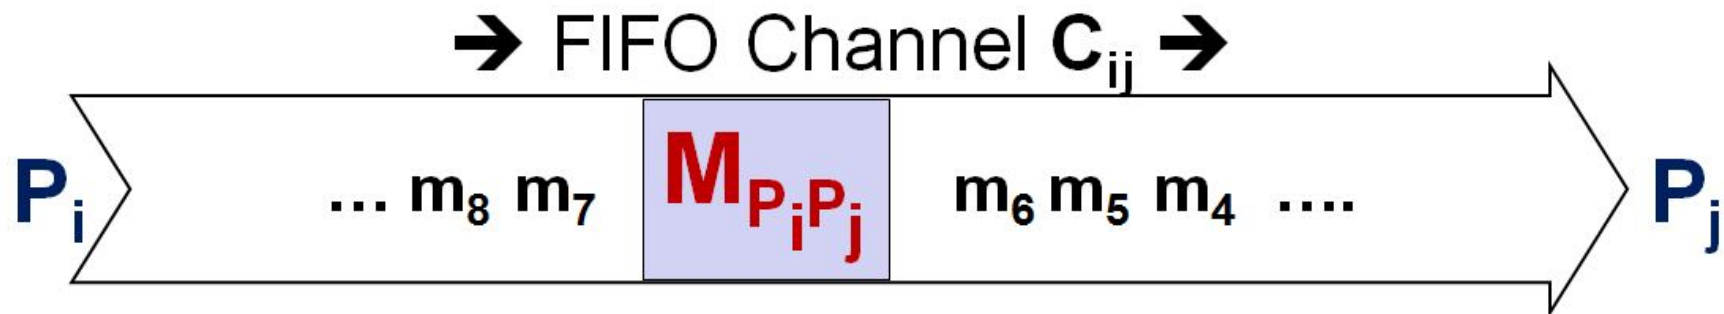
Record GS ; Update GS after waiting t_{max}
 \implies strong consistent GS' is achieved easily.

Chandy-Lamport Snapshot Algorithm – 1

► **Precondition:** lossless FIFO channels

► **Idea:** Isolate Phase^k; Record states; Phase^{k+1}

1. Initiate state recording **locally** in any process P_i spontaneously
2. **Propagate** global recording via **marker** msg on all channels
 FIFO eases isolation between different phases by distinguishing among messages sent **before/after** local state recording.
 \implies Processes are allowed to proceed including snd/rcv actions.



3. Each P_i stores all msgs received after most recent recording LP_i
4. At the end of state recording \longrightarrow send state to Coordinator

Chandy-Lamport Snapshot Algorithm - 2

- **Start Snapshot Algorithm:** in an arbitrary process P_i

record(LP_i); (stores LP_i locally)

FORALL $C_{i,j}$ DO snd($M_{i,j}$) OD; (propagate execution)
- **Reaction** on a rcv($M_{i,j}$) in process P_j (react at all times)

IF (NOT recorded(LP_j)) THEN

STATE($C_{i,j}$) := empty; record(LP_j); (store locally)

FORALL $C_{j,k}$ DO snd($M_{j,k}$) OD; (propagate)

ELSE (update local state by all $msg_{i,j}$ received after local recording)

STATE($C_{i,j}$) := STATE($C_{i,j}$) \cup

$\{msg_{i,j} \mid time(LP_j) < time(rcv(msg_{i,j})) < time(rcv(M_{i,j})) \}$;

FI;
- **Combine to global state:** all processes send LP_i to coordinator after all marker have arrived in P_i (counter for $|\{C_{i,j}\}|$)

Observations – Algorithm Properties

- ▷ **Result:** consistent global state
 - no:** strong consistent state, **but:** states of channels are known
- ◁ **complete global state** $\implies \exists P_s$ connected to *all* processes in PS
- ▶ **Process Structure:** Initiation
 - **Diffusion** phase \approx when first marker arrives in a process P_j
forward marker $M_{j,k}$ to all processes P_k directly connected to P_j
 - **Contraction** phase \approx update local state for 'late' messages
no more forwarding of markers
 - **Collection** phase \approx send local state(s) to coordinator process
- ▶ **Distinguishing 'runs'** for different 'spontaneous' initiations:
 1. Extend marker $M_{j,k}$ by initiator process number P_i : $M_{i,j,k}$
 P_i acts as coordinator \implies decentralized, independent 'runs'
 2. Additional counter in P_i discriminates different runs by P_i

important
w.r.t.
termination

Properties of Snapshot Algorithm – 1

1. **Correctness:** Resulting state is consistent \implies **no Orphans**

Proof: (indirect)

Assumption: global state holds (at least) one **Orphan** message

$\implies \exists msg_{i,j}$ s.t. $\text{snd}(msg_{i,j}) \notin LP_i \wedge \text{rcv}(msg_{i,j}) \in LP_j$

$\implies P_i$ sends msg **after** recording local state LP_i and

P_j has received msg **before** recording local state LP_j

\implies **Contradiction:** Processes do not respect protocol:

(a) $msg_{i,j}$ received **before** marker $M_{i,j}$ in $P_j \implies$

i. execution order: $\text{snd}(msg_{i,j}); \text{snd}(M_{i,j})$ in P_i **or** (illegal)

ii. Channel $C_{i,j}$ is not FIFO (precondition not met)

(b) $msg_{i,j}$ received **after** marker $M_{i,j}$ in P_j

execution order: $\text{rcv}(M_{i,j}); \text{rcv}(msg_{i,j}); LP_j$ in P_j (illegal)

\implies **No Orphans iff algorithm is executed as programmed.**

Properties of Snapshot Algorithm – 2

2. **Termination:** for a single, distinguishable run

Precondition: isolated steps in P_j terminate; finite transfer times

- Diffusion phase initiated by P_i
at most $|PS|*(|PS|-1)$ markers (maximum connectivity) as each P_j store and propagates exactly once (THEN) c.f. pg. VI-63
- Contraction phase ends after a maximum of $(|PS|-1)$ received markers at each process (full connectivity)
- Collection phase: exactly $(|PS|-1)$ messages LP_j go to P_i .

3. **Costs:** local memory, compute time and extra messages

- $|PS|$ local states LP_i and $|PS|*(|PS|-1)$ channels
- $|PS|*(|PS|-1)$ markers and $(|PS|-1)$ local states

Option: Record $B \subset PS$ of 'important' processes only,
e.g., *Server* for data bases, long-running computations

Importance of Message Channels for Snapshots

- ◀ Initiator does **not** reach all processes \implies **incomplete** snapshots
direct connection or connectivity via transitive hull of **msg** channels
- ◀ explicit control messages needed to avoid incomplete snapshots
otherwise: input data and actual process run may not contact all
processes \implies not all processes will get markers.

Tradeoff: *Additional control messages vs. quality of result*

- ◀ **Channel properties** are essential:

1. No FIFO property \implies simulate message ordering
2. No *piggy-backing* of messages or extra messages permitted?

c.f.
VI.2

Without (1) **and** (2) \implies Algorithm has to **freeze entire system**
until snapshot is recorded. (*How?*)

Taylor
1989

at least: prevent processes from sending messages during recording!

Note: Determining *TRANSIT* is easy for FIFO channels with
known maximum transfer times, otherwise hard to get.

c.f.
pg.
VI-60

Global States and Causality


- **Cut** of process system $PS \approx$ *exactly one event for each* $P_i \in PS$
 - \implies Recording $LP_i \forall P_i \in PS$ results in a *cut* of PS
 - Cut is consistent* \iff all events are pairwise **concurrent**
 - i.e. are not ordered w.r.t. *happened-before* relation $\xrightarrow{\sqsubseteq}$
 - \implies **consistent global state is also a consistent cut**
- Orphan:** rcv without snd implies causal chain through channels

- **Alternative Formalization:** (*Petrinet/process theories*) c.f. pg. VI-69
 1. $e_{ik} \in P_i$ subsumes its **history**(e_{ik}) $:= e_{i1}e_{i2} \dots e_{ik}$, i.e., is identified with the prefix of the causal chain that leads to e_{ik}
 2. **Cut**(PS) $:= \bigcup_{i \in [1:|PS|]} \text{history}(e_{ik_i})$ (1 chain for each process)
 3. **Front** $:= \{e_{1k_1}, e_{2k_2}, \dots, e_{nk_n}\}$ most recent events of processes
 4. **Cut**(PS) *is consistent* \iff

$$(b \in \text{Cut}(PS) \wedge a \xrightarrow{\sqsubseteq} b) \implies a \in \text{Cut}(PS)$$
- 2. ensures local order in P_i ; **msgs** are the 'critical' events

Example: Synchronic Distance among Processes

c.f.
Colouris
et al.:
Distributed
Systems
Fig.10.15



- 'Lattice' of all permitted, combined states for P_1 and P_2
- more de-coupling, e.g., S_{25} not permitted w.r.t. consistency

End
VI.4

VI.5 Detecting Global System States

Global Consistent View on a DS is always a challenge!

► **Non-volatile, permanent States:** Detection not time-critical

- **System Termination**
- **Deadlocks**

VI.5.1

VI.5.2

◄ **Volatile States:** How to obtain a consistent view?

VI.4

- local changes may occur spontaneously
- detection algorithm may change the system behaviour
- what about messages in *TRANSIT*?

Additional Dimension: *State reachable in every run of the system?*

► Situation in the 'running' PS at hand \implies Debugging


◄ **Property:** for all PS when running a program \implies Proof

- * *Safety*-Properties: nothing bad will ever happen
- * *Liveness*-Properties: something wanted is eventually reachable

VI.5.1 Detecting System Termination

System Model: $PS = \{P_1, \dots, P_n\}$

1. Message channel C_{ij} between each process pair (P_i, P_j)
channels are 'robust', i.e. no messages are lost
2. Every $P_i \in PS$ is always in one of the two states { idle, active }:
 - (a) active \implies Process may snd, rcv or compute
 - (b) idle \implies rcv of messages only (**NO** snd!)
3. All State Changes permitted but:
no spontaneous change to **active** based on internal reasons only



Why Termination? always a problem in de-centralized algorithms
 \implies do not 'mix' termination detection with application

Basic Algorithmic Idea

- ▶ **Initial State:** Every process is idle, no messages
 - $\forall P_i \in PS: STATE(P_i) = \text{idle}$
 - $\forall P_i, P_j \in PS: TRANSIT(P_i, P_j) = \emptyset$
- ▶ **Final State:** the same as initial state
- ▶ **Method:** unique **Monitor Agent** oversees all activities
 - \exists **exactly one** $P_{mon} \in PS$ where $P.\text{monitor}() = \text{true}$
 - **Weights** represent the level of work a process has to do
 1. $\forall P_i \in PS$ holds $\text{Weight}(P_i) \geq 0$
 2. **Algorithmic Invariant:**

$$\sum_{P_i \in PS} \text{Weight}(P_i) = 1$$
 - P_{mon} starts with the initial job and the $\text{Weight}(P_{mon}) = 1$
 - Processes get part of the job along with part of the Weight
 - If a job is done, the corresponding weight is sent back to P_{mon}
- Result:** $(\text{Weight}(P_{mon}) = 1) \implies$ System is terminated

VI-60

Termination Detection Algorithm (Huang)

```

STATE := idle; permitted_splits := S; W := 0;                                (0)
/* fixed number S avoids non-terminating distribution */
DO (non-deterministic choice among matching alternatives)
  rcv( $P_{from}$ , WORK, Weight)  $\longrightarrow$  W := W+Weight; /* may happen anytime */ (1)
  IF (STATE==idle) THEN STATE := active; FI;
  rcv( $P_{from}$ , CTRL, Weight)  $\longrightarrow$  W := W+Weight; /* Monitor  $P_{mon}$  only */ (2)
  IF (W==1) THEN snd( $P_{out}$ ,TERM); STATE := idle; FI;
  (STATE == active)  $\longrightarrow$  /* at least once, at most S times work distribution */ (3)
  IF (permitted_splits > 0) THEN permitted_splits := permitted_splits-1;
    <  $W_1, W_2$  > := split(W); /* where  $W_1, W_2 > MIN \wedge W_1 + W_2 = W$  */
    W :=  $W_1$ ; P := choose( $PS \setminus \{self\}$ ); snd(P,WORK, $W_2$ ); FI;
  (STATE == active)  $\longrightarrow$  /* Do Real Work */ (4)
  perform(WORK);
  IF (NOT(monitor())) THEN snd( $P_{mon}$ ,CTRL,W); W := 0; STATE := idle; FI; (4a)
OD;

```

Assumption: initial msg ($P_{out}, Work, 1$) is received by P_{mon} in (1)

Outline of Correctness and Termination Proof

1. P_{out} initializes algorithm with weight 1 and activates P_{mon} in (1)
2. `split` is a loss-less division \implies no weights are lost
3. **Invariant:** $W_{mon} + W_{active} + W_{work} + W_{ctrl} = 1$ (after init)
 - $W_{active} \approx$ Sum of weights from all P_i processes except P_{mon}
 - $W_{work}, W_{ctrl} \approx$ Sum of weights of corresponding messages
4. $P_i \in PS \setminus \{P_{mon}\}$ active $\iff W_{P_i} > 0$
 - initially: $W = 0$ and `STATE` = `idle`
 - (`idle` \mapsto `active`): only in step (1) adds `Weight` > 0
 - (`active` \mapsto `idle`): only in step (4a) and $W := 0$ afterwards
5. Algorithm terminates $\iff P_{mon}$ sends `TERM` to P_{out}
 Invariant 3. where $W_{active} + W_{work} = 0$; $W_{ctrl} = 0$ after `TERM`

Termination: perform/msg-transfer times **finite** (Assumption)

- Only finite number of `split` steps in (3) due to ($MIN > 0$)
- No infinite delegation cycles in (3;1) due to limit S

VI.5.2 Distributed Deadlock Detection

Resources: *'Everything that processes compete for'*

see
also
REST

- ▶ devices scheduled/data provided by the operating system
Transparency \implies local as well as remote devices, data, ...
- ▶ Synchronization: Access to critical sections
- ▶ Messages/RPCs: blocking wait for a message or reply

VI.2.2

Resources also an issue in traditional OS:

- OS Usage Protocol: Request – Wait – Hold – Free
- Deadlock problem if scheduling is bad or too optimistic
Prevention vs. Avoidance, Detect and Resolve, Ignore

[**A. Tanenbaum, 1995**]:

Deadlocks in distributed systems are similar to deadlocks in single-processor-systems, only worse.

Goal: Efficient and Fair Resource Management

- ▶ **Global View** about state of all resources **required**
 - **centralized** is easy to control but an architectural **bottleneck**
⇒ all requests go through a **central Server/Manager**
 - **de-centralized** ⇒ problems w.r.t. **consistency of view**
e.g., resource states change during message transfer times
- ◀ **Problem for both models:**
messages (requests, ..., states) take too long, get lost, re-ordered
⇒ **General Problem requires handling of volatile states!**

c.f.
pg.
VI-70

Resource Model: (simplified for our algorithms)

- **non-consumable**, i.e. re-usable resources with **exclusive access**
- **no** multiple equivalent instances of resources (type vs. instance)

VI-78

System Models for Processes PS & Resources RS

System State assignment of internal/external resources to processes that are either internal or on remote nodes

1. Resource-Allocation Graph $(PS \cup RS, E_{req} \cup E_{assign})$

- $E_{req} \subseteq PS \times RS := \{(P_i, R_j) \mid P_i \text{ waits for } R_j\}$ (**requests**)
- $E_{assign} \subseteq RS \times PS := \{(R_j, P_i) \mid P_i \text{ holds } R_j\}$ (**assigned**)

Detailed: where are resources and who waits for which resource

2. Wait-For-Graph (PS, E) where $E \subseteq PS \times PS$ and $(P_i, P_j) \in E \iff \exists \text{ Resource } R \text{ s.t. } P_i \text{ is waiting for } R$ and $P_j \text{ holds } R$

Abstraction: hold \approx unambiguously

wait \approx maybe more than one process waits?

Deadlock $D \subseteq PS \iff (D, E \cap D \times D)$ contains at least one cycle

Note: Cycle implies deadlock iff there is only one instance/resource

\implies *Deadlock Detection is done via Cycle Detection*

Problem: RAG vs. WFG – Multiple Instances

Legend:

Circle $\in PS$

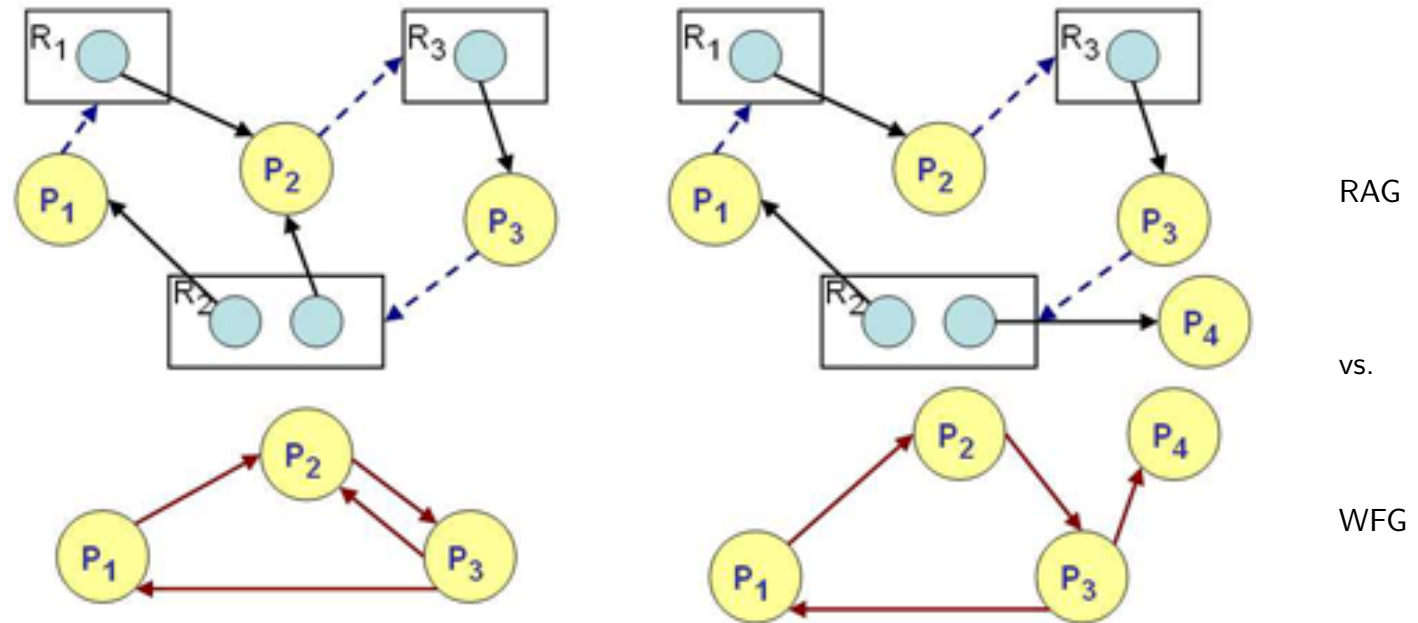
Square $\in RS$

Circle in Square

$\approx RS$ -Instance

Edge (PS, RS) request

Edge (RS, PS) assign



WFG describes Wait-Relation, but no detailed cause w.r.t. resources

► more compact \implies easier to find cycles

◄ cycles may imply **false deadlocks**

▷ **left:** conflict establishes cycle and a deadlock

◁ **right:** conflict establishes cycle but *no deadlock* (P_4 may end)

Reason: Multiple resource instances imply **OR**-Relation for wait

Distributed Deadlock-Handling: Detect & Resolve

◀ **Centralized Solution:** RAG/WFG in a global control process

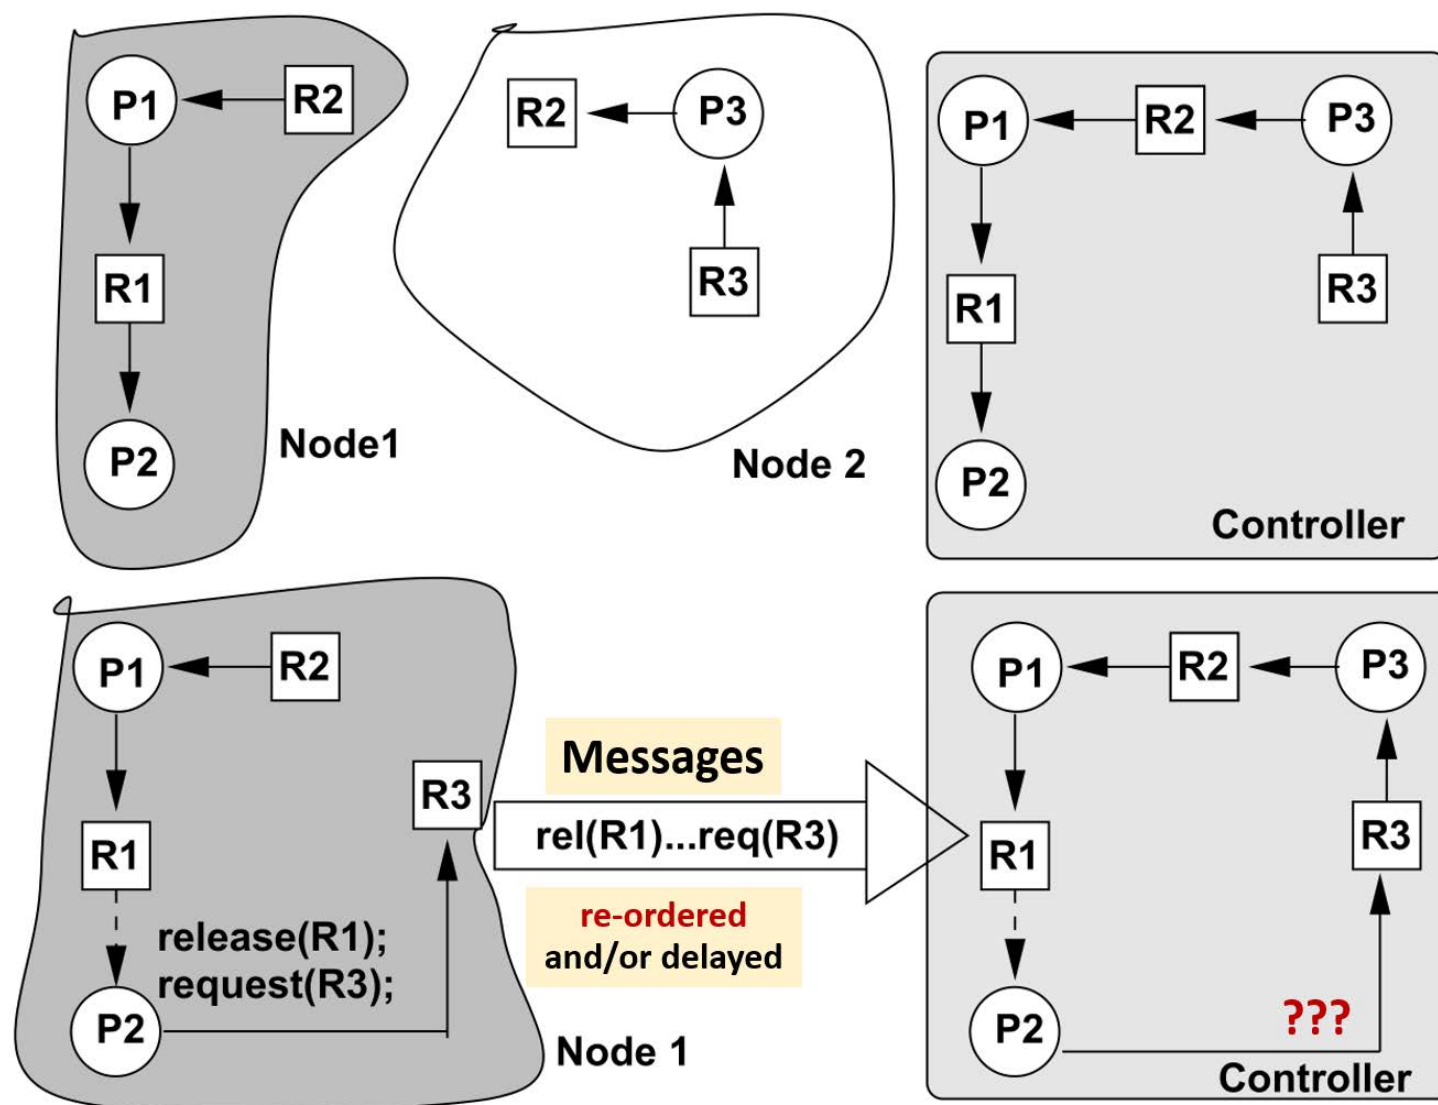
▶ **Adopted Centralized Solution:**

1. **Local** Resource-Allocation-Graph in each process VI-80
2. **Controller** combines local RAGs for detection; updates **either**:
 - (a) $P_i \in PS$ send always all changes to Controller push
 - (b) $P_i \in PS$ send local RAGs in fixed intervals push
 - (c) Controller requests local RAGs on demand pull

Problem: **false deadlocks** due to message timings \implies

- ▷ ordered msg channels or request updated infos (timestamps)
- ▷ Controller holds 2 global RAGs and computes AND for all edges
 \implies 2-Phase Ho-Ramamoorthy (1982)

▶ **Distributed Solution:** distributed construction of Wait-For-Graph VI-82
out of local WFGs and **Edge Chasing Algorithm**
 \implies Chandy-Misra-Haas (1983)



False Deadlock identifiable: P_2 frees R_1 ; Termination Order $P_1 \longrightarrow P_3 \longrightarrow P_2$

Distributed Edge-Chasing–Algorithm

- ▶ Processes: request of multiple resources in a single request
 \implies Process may wait for more than one other process
- ▶ Distinction: Waiting for **local** vs. **remote** resources
- ◀ No Controller \implies No global knowledge in a single process

Algorithmic Idea: How to detect cycles?

If process waits some defined time for a remote resource (blocked)
 \implies Deadlock is 'suspected' \implies Msg exchange starts

- ▶ **Probe Msg:** $\langle i, j, k \rangle \approx \langle \text{Initiator, Sender, Destination} \rangle$
- ▶ **Procedure:**


- **Initiator** P_i sends Msg to all processes he **waits** for
- also **waiting** destinations send msg to all those they wait for
- If process P_i receives a message, it initiated \implies **Deadlock**

- ▶ **Optimization:** local waits are handled by local graph, not msgs
 msgs store route as a hint for deadlock resolution

VI-82

Expl.: Edge-Chasing using three distributed Nodes

- local messages are avoided by looking into local Resource Graph
- msgs between nodes: explicit message-passing



Advantage: Construction of 'global knowledge' On Demand

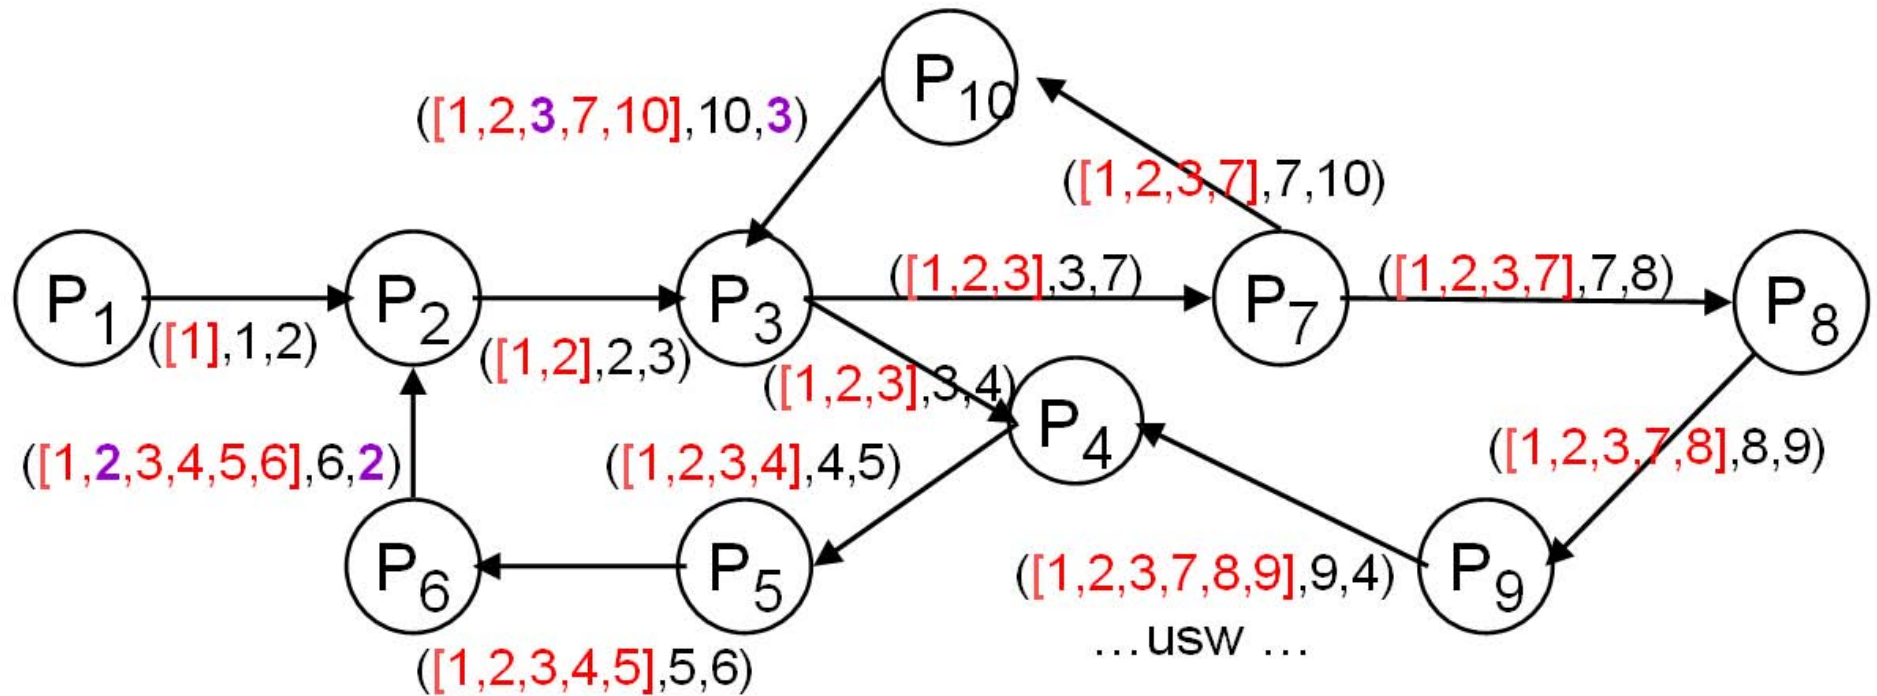
- ▶ Times for messaging not critical: processes are blocked anyway
- ◀ Waiting processes have to react to inquiry ([watchdog](#))
- ▶ single local controller with local RAG per node

Remark: *lots of different algorithms in literature*

Thread

Shin-
gal
et
al.

Expl.: Optimized Edge-Chasing stores Routes



Method: Propagate list of **all** nodes visited

Each node check whether it is the list

- Deadlocks are found independent of initiating process
- ◄ If multiple deadlocks are found \implies handling more complicated

Resolution: Preempt & Withdraw Resources

- ◀ **Problem:** preempted process loses work already done

recovery-points \neq process start \implies costly

VI.4

Acceptable iff recovery mechanisms are required anyway

- ▶ **Distributed Databases:** Transaction Management and Logging

- Precondition: unique time stamps used for prioritization

$t_{P_1} < t_{P_2} \implies P_1$ older than $P_2 \implies P_1$ loses more work

\implies **Older Process will not be terminated!**

- **How to avoid cyclic 'killings'?**

- ◁ **Wait-Die:** wait only for **higher** time stamps

$P_{old} \longrightarrow P_{young} \approx \text{wait} \mid P_{young} \longrightarrow P_{old} \approx \text{kill}(P_{young})$

Problem: $(\text{kill}; \text{restart})^+ \implies$ **thrashing** effect

- ▷ **Wound-Wait:** wait only for **lesser** time stamps

$P_{old} \longrightarrow P_{young} \approx \text{kill}(P_{young}) \mid P_{young} \longrightarrow P_{old} \approx \text{wait}$

Advantage: kill; restart; wait much less overhead

End
VI.5

Remark: *Lots of literature in distributed OS and DBS.*

VI.6 Distributed Coordination

- ▶ Arbitrarily distributed application system with *distributed*
 - * *control* to ensure robustness without bottlenecks
 - * *data* due to replication-based transparency
 - * *compute load* due to efficiency and feasibility reasons
- ◀ Basic level of agreement and consensus essential for, e.g.,
 - * mutual exclusion and consistency mechanisms for replicated data
 - * centralized/de-centralized organization of common decisions

Typical coordination problems in distributed systems

1. **Election**: Determine new unique coordinator after crashes
2. *Agreement* about common 'global' state in the context of errors
3. *Commitment* about executing collective actions, e.g. *transactions*

additi-
onal
MSc
topics

Note: Dynamic process systems/groups with ever changing process numbers are especially 'hard', e.g., mobile or P2P systems.

VI.6.1 Leader Election

Motivation: Leader Election is the basis for

- ▶ distributed algorithms that use some kind of centralized organization, e.g., mutual-exclusion, resource handling, deadlock or termination detection
- ▶ fault-tolerance in these algorithms by allowing for de-centralized replacement of crashed coordinators, servers etc.

⇒ **Requirements for any Election Algorithm:**

1. may be initiated by **any** process in the system
2. may be initiated **in parallel** by different processes

Reason: a priori unknown when and where a process crashes

3. terminates always appointing **exactly one Leader**
i.e, all processes know their state w.r.t. election: $\{ \text{leader}, \text{lost} \}$

System Model for Leader Election

- ▶ **Lossless** message transfer
- ▶ Each $P_i \in PS = \{P_1, \dots, P_n\}$ knows about **all 'names'** in PS realistic ?
- ▷ All processes $P_i \in PS$ are **a priori equal** *
 \implies **Each process** is able to act as a coordinator
- ▷ **Priorities** are defined by arbitrary **total, strict order** $<$ on PS
 e.g., totally ordered process indices $P_1 < P_2 < \dots < P_n$
 $\forall M \in \wp(PS)$ is $P_i \in M$ assigned the highest priority $:\iff$

$$i = MAX(\{j \mid P_j \in M\}) \quad *$$

At any time, PS is partitioned into two process sets:

1. $UP(PS) \approx$ active processes
2. $DOWN(PS) \approx$ crashed processes

State changes are **spontaneous and not coordinated**

Caution: Even change from *crashed* \longrightarrow *active* is **not** predictable

Additional Challenges

- ◀ **Distinction:** crashed nodes vs. lost messages \implies
lossless communication with **maximum response time** t_{max} known
(msg transfer plus time to react, e.g., via 'watchdog' thread)
- ◀ **Elections may be initialized in parallel**
set of nodes waiting for a response governed by a **timeout**
e.g., access to a **critical section**: missing WAIT/OK messages c.f.
III-66
- ◀ **Formerly Leader is re-started** and becomes active again
Who is assumed to stay/become coordinator in this case?

System structure and fault tolerance

1. **complete** reachability \implies *broadcast*
2. **partial** reachability \implies special protocols needed
Example: Algorithms based on ring structures

Note: Partitioning due to channel crashes is most critical

Challenge: Determine a **single**, unique coordinator, s.t.
all $P_i \in UP(PS)$ **agree in**
the name of the (new) coordinator process

Typical Algorithms: *Bully*-Algorithm
Ring-Algorithm

c.f.
pg.
VI-93

Bully Algorithm (Garcia-Molina 1982)

- ▶ **Precondition:** Each active process is able to reach **all** others
- ▶ **Idea:** Prioritize the process with the highest process index
- ▶ **Method:** timeout in a process \implies
 - initiate re-election for all processes with higher indices
 - stop re-elections from processes with lower indices
 - single process that has highest index \implies is elected

Note: start of multiple re-elections in parallel possible
after re-start, the 'former' coordinator is favored to become
the new coordinator again ('bullies' the current coordinator)

1. **Time-out** of a *msg* to (last known) coordinator \longrightarrow GOTO 2.

```

FORALL  $\{P_j \in PS \mid j > i\}$  DO snd( $P_j$ ,VOTE) OD;      /* higher up? */
ELECTION := true; RESPONSE :=  $\emptyset$ ; Coord := undef;
Wait( $T$ ) for EVENTS of  $\{3, 5\}$ ;                          (timeout  $T$ )
IF (RESPONSE ==  $\emptyset$ )                                /* all higher procs down */
    THEN FORALL  $\{P_j \in PS \mid j < i\}$  DO snd( $P_j$ ,OK) OD;      /* self */
    ELSE Wait( $T'$ ) for EVENTS OF  $\{4, 5\}$ ;          /* higher elect */
        IF (Coord == undef) THEN 2. FI
ELECTION := false;

```

$$3. \text{rcv}(P_j, \text{WAIT}) \longrightarrow \text{RESPONSE} := \text{RESPONSE} \cup \{P_j\};$$

4. $\text{rcv}(P_j, 0K) \longrightarrow \text{Coord} := P_j;$ /* new coordinator? */

```

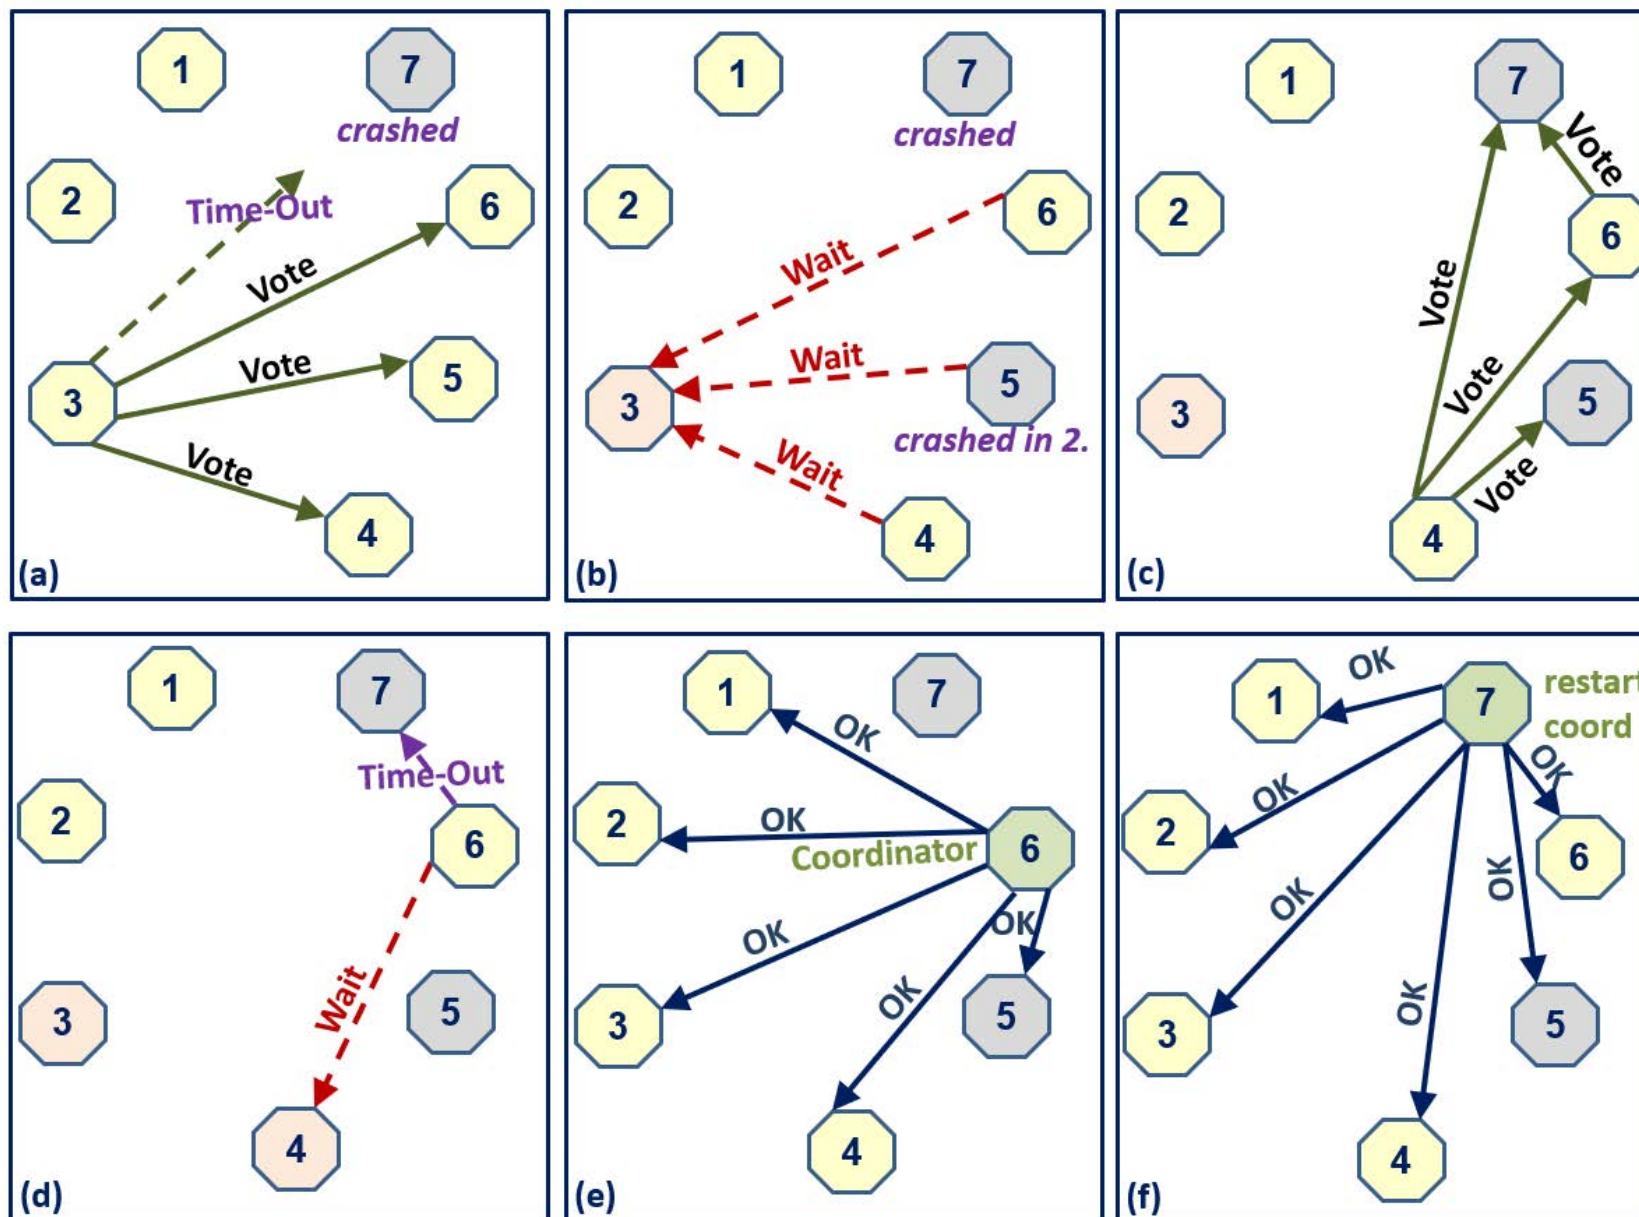
5. rcv( $P_j$ , VOTE)  $\longrightarrow$  snd( $P_j$ , WAIT);      /* stop 'lower' election */
    IF NOT(ELECTION) THEN 2. FI;                      /* vote */

```

6. **Recover** (from crash) \longrightarrow GOTO 2.

Note: Process crashes after WAIT \implies never 4.: re-start election

Example: Bully Algorithm



Assessment of Bully Algorithm

- ▶ **Nomen est Omen**; very simple structure
- ▶ works despite multiple node crashes during election
- ▶ parallel elections cause no problems and will be stopped 'early'
 \implies processes with higher process indices determine overall result
- ▶ a single run terminates always due to finite number of messages

- ◀ Lots of crashes \implies triggers incessantly re-elections!
 \implies Determination of **timeouts** is critical for success

- ◀ **Costs** $\mathcal{O}(|PS|^2)$ messages (*worst-case*):
 UP = PS; P_n crashes;
 P_1, \dots, P_{n-1} observe crash in parallel and start $(n - 1)$ elections;
 P_n is restarted again
 $\sum_0^{n-1} \text{VOTE} + \sum_0^{n-1} \text{WAIT} + (n - 1) \text{OK messages}$

Conclusion: rather costly w.r.t. number of messages exchanged

Ring Algorithm

1. **Naive** version using direct neighbors in a ring
2. **Efficient** version collecting information in a logical tree structure

Ring-Algorithm (LeLann 1977)

- ▶ **Precondition:** uni-directional ring for communication among active processes $P_1 \mapsto P_2 \dots \mapsto P_n \mapsto P_1$
 $P_i.next()$ determines current next **active** neighbor $\in UP$
 \implies *strong precondition used here!*
- ▶ **Idea:** Use maximum index in **list of processes** w.r.t. ring direction
- ▶ **Method:** timeout in process $P \implies$
initiate re-election; start with singleton list $[P]$
handover and extend list
after complete circle \implies propagate coordinator list

Note: start of multiple re-elections in parallel possible
restart of a former coordinator initiates new election

LeLann Ring Algorithm


Each active process P_i reacts to the following events:

1. **Time-out** of a msg to (last known) coordinator \longrightarrow GOTO 2.
2. **Initiate Election:** $\text{snd}(P_i.\text{next}(), \text{VOTE}[i])$;
3. $\text{rcv}(P_j, \text{VOTE}[List]) \longrightarrow$ /* transform */
 IF $i \in List$ THEN $\text{snd}(P_i.\text{next}(), \text{COORD}[List])$;
 ELSE $\text{snd}(P_i.\text{next}(), \text{VOTE}[i : List])$;
 FI; /* extend */
4. $\text{rcv}(P_j, \text{COORD}[List]) \longrightarrow$ /* term/propagate */
 Coord := MAX([List])
 IF $i \neq List.\text{front}()$ THEN $\text{snd}(P_i.\text{next}(), \text{COORD}[List])$;
5. **recover** after crash \longrightarrow GOTO 2. /* no privileges */

Note: $VOTE \approx$ ask around; $COORD \approx$ propagate result

Termination: initiating process aborts further propagation

Example: LeLann Ring Algorithm



Assessment of LeLann Algorithm

- ▶ simple structure
- ▶ parallel elections cause no problems and achieve common result MAX
- ◀ **Costs:** each election requires $2 * (|PS| - 1)$ messages
 worst-case $(|PS| - 1)$ parallel election procedures
 $(|PS| - 1) * (2 * |PS| - 1)$ messages $\implies \mathcal{O}(|PS|^2)$

Conclusion: rather costly and requires optimizations

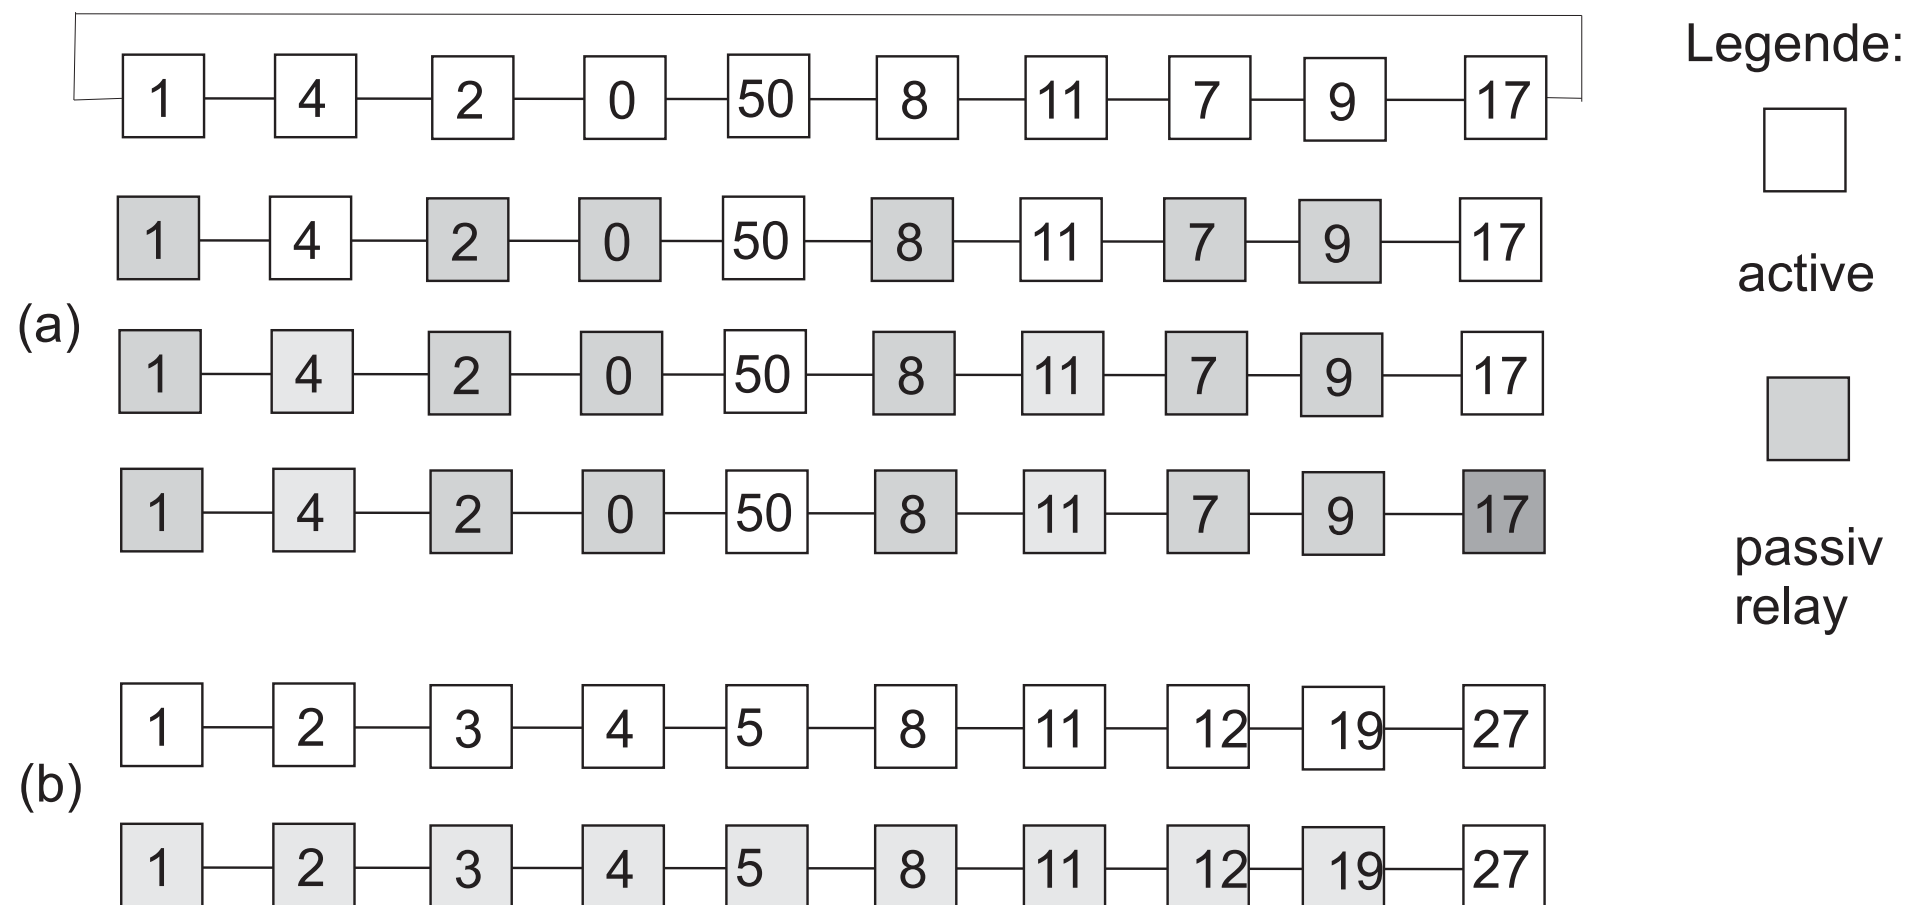
Worst-case Optimization: (Peterson 1982)

- tree-like reduction of ring structure for parallel elections
- **active** \approx participate in election / **relay** \approx transfer messages only
- each step cuts number of **active** processes in half (at least) via
 - **bi-directional** ring: $\text{MAX}(\{P_{i \ominus 1}, P_i, P_{i \oplus 1}\})$ (left/right)
 - **uni-directional** ring: $\text{MAX}(\{P_{i \ominus 2}, P_{i \ominus 1}, P_i\})$ (2 predecessors)

ACM
ToPLaS,
(10)
1982

c.f.
pg.
VI-97

Example: Peterson **Ring-Election** Algorithm



VI.6.2 Agreement Protocols – Outlook

Simplified Process Model

- PS with n Processes (nodes)
- **robust, loss-less direct communication** $\forall (P_i, P_j) \in PS^2$
 \implies no special handling based on node where error occurs

Reason: Network Partitioning in central node, e.g., star network

Error in **name server** vs. user node

No subsequent errors due to transitive message routing

ERROR (Fault): *Deviation of expected system behavior* in

- ◀ **General** systems, i.e., synchronous as well as asynchronous:
no or unexpected reaction of nodes or msg channels
- ◀ **synchronous** systems additionally via **timing** errors, e.g.,
clock drift exceeds tolerable limit; msg transfer times
time for computing or replies exceeds limit

VI-100

c.f.
I-33

Why is Agreement important in DS?

- ◀ **Quality of Service** assumptions are not met
Expl.: E-commerce, online orders ...
guaranteed delivery time exceeded; order not processed
- ◀ **Distributed Algorithms fail** due to preconditions not met
⇒ distributed infrastructure no longer robust
Expl.: Algorithms waiting for msgs will block
Bully-Algorithm using faked process indices

⇒ **Fault tolerance is important**

- ▶ **Replication** of active and passive components
- ▶ Saving states as recovery points; Logging (transactions)
- ▶ **Reduce impact of faulty components**
 - * Determine which parts are faulty and which are ok
 - * Secure functionality of system parts working correctly

chap.
VII
VI.4

Errors and Faults: Causes and Classes

- **Connections:** lost, duplicated, corrupted msgs
- **Nodes:** different levels of impact for errors, esp.
 1. **Fail stop:** P_i **identifiable** permanent **down**
 2. **Crash fault:** P_i **not identifiable** permanent **down**
 \implies Process does not send or reply in the future
 3. **Send Omission fault:** not all msgs (to all destinations) are sent broad-
cast
 4. **Receive Omission fault:** not all msgs arriving are accepted
Impact: content loss plus blocking in synchronous interaction
 5. **Byzantine fault:** Unpredictable behavior of nodes !! worst
case
sometimes correct, sometimes faulty behavior
omits/**manipulates** some msgs, even generates unexpected msgs

AGREEMENT: Make sure that **all non-faulty nodes** get always the correct information

\implies *Prohibit or Reduce the effects of faulty nodes.*

Agreement Levels for Byzantine Faults

1. Byzantine Agreement (single-source broadcast)

A **single** P_i propagates a fixed value v_i and **all** non-faulty processes agree on a **single value** V (If P_i is not faulty $\implies V = v_i$)

2. General Consensus (multiple-source broadcast)

Each P_i propagates **it's own** fixed value v_i (values may differ) and **all** non-faulty processes agree on a **single value** V
(If all v_i are the same in all non-faulty processes $\implies V = v_i$)

arbitrary
?

3. Interactive Consistency (multiple-source broadcast)


Each P_i propagates **it's own** fixed value v_i (values may differ) and **all** non-faulty processes agree on a **Vector** (V_1, V_2, \dots, V_n) of values
(If P_i is not faulty \implies entry V_i of the vector V is v_i)

see
Colou-
ris
11.5

Note: Implementing 1. for all processes implies 3.

Implementing 3. plus a global majority function implies 2.

Byzantine Generals – 'History?'



From cs4410 fall 08 lecture

Fig.: gauthamzz.github.io/tendermint.html#byzantine-fault-tolerant

Byzantine Agreement - Basic Situation

- ▶ A single P_i wants to establish a consensus about value v_i in PS
- ▶ Arbitrary node failures are 'expected'; no messages lost VI-100
- ▶ All processes use the same symmetrical algorithm
- ◀ Messages may be manipulated/faked
- ▷ **No** use of authorization or signatures VI-108

Algorithm: Exchange values to guarantee **forming a majority**
Problem:

- ◀ Nodes may manipulate msgs before forwarding
 - ⇒ multiple interaction (rounds) used to detect faked msgs
 - ⇒ message exchange is very costly
 - ◀ **Only** $m \leq \frac{n-1}{3}$ (*ceiling*) of **faulty processes** in n nodes are tolerable, i.e., $n = 3m + 1$ processes may compensate for m faults VI-104
- Example: 1 out of 4 is ok; 1 out of 3 not;
2 out of 7 processes etc.




Example: 3 Processes with 1 Faulty Process

P_2 : P_1, P_2 ok, P_3 fault $\implies P_2$ has to choose P_1 value

left

P_2, P_3 ok, P_1 fault $\implies P_2$ has to choose P_1 value (same algorithm)

center



P_3 : P_1, P_3 ok, P_2 fault $\implies P_3$ has to choose value of P_1

right

P_2, P_3 ok, P_1 fault $\implies P_2$ has to choose value of P_1 (same alg.)

center

\implies **No Agreement** among P_2 and P_3 if P_1 is faulty

Generalization: ($m > 1$) reduces to 1-of-3 problem by *contradiction*
 \implies general algorithm would work also for ($m = 1$)

Oral-Message-Algorithm: Lamport/Shostak/Pease

ACM
ToPLaS
1982

Method: Recursive Exchange as Remote Procedure Calls

► **Arbitrary Call** $\text{om}(m, \text{PROCS}, p)$

$m \approx \text{Max-Faults}$; $\text{PROCS} \subseteq PS$ participate; $p \approx \text{starts call}$

$\text{om}(m, \text{PROCS}, p) ::= \text{PROCS} := \text{PROCS} \setminus \{p\};$
 FORALL $P \in \text{PROCS}$ DO $\text{snd}(P, \text{val}_p)$ OD; (1)

IF ($m > 0$) THEN
 FORALL $P \in \text{PROCS}$ DO $\text{RPC}(P, \text{om}(m-1, \text{PROCS}, P))$ OD; (2)

 FORALL $P \in \text{PROCS}$ DO $\text{RPC}(P, \text{majority}(P, \text{PROCS}))$ OD; (3)

FI;

$\text{majority}(P, \text{PROCS}) ::= \text{val}_p := \text{Majority value from all } P' \in \text{PROCS};$

► $P \in PS$ **reacts:** after $\text{rcv}(p, \text{val}_p)$ start own RPC call

► **initial Call:** $\text{om}(m, PS, P_{init})$ of Initiator P_{init} using $\text{val}_{P_{init}}$

Costs: $\text{om}(m, PS, p) \approx |PS|-1 \text{ RPCs } \text{om}(m-1, PS \setminus \{p\}, P) \dots$

Expl.: $m = 1$; 4 processes $PS = \{P_0, P_1, P_2, P_3\}$: $\text{om}(1, PS, P_0)$

1. Case: Initiator P_0 is non-faulty

(a) Let $\text{val}_{P_0} = 0$

P_0 sends 0 to $\{P_1, P_2, P_3\}$

$\text{val}_{P_1} = \text{val}_{P_3} = 0$


(b) in P_1 $\text{om}(0, \{P_1, P_2, P_3\}, P_1)$

in P_2 $\text{om}(0, \{P_1, P_2, P_3\}, P_2)$

in P_3 $\text{om}(0, \{P_1, P_2, P_3\}, P_3)$

forward correct in P_1 and P_3

forward faked in P_2



(c) majority

in P_1 using $\langle 0, 1, 0 \rangle \implies 0$

in P_2 using $\langle -, -, - \rangle ??$

in P_3 using $\langle 0, 1, 0 \rangle \implies 0$


$\implies \{P_0, P_1, P_3\}$ agree on the **same value 0**

$/^* \text{ values from } \langle P_1, P_2, P_3 \rangle ^*/$

Expl.: $m = 1$; 4 processes $PS = \{P_0, P_1, P_2, P_3\}$: $\text{om}(1, PS, P_0)$

2. Case: Initiator P_0 itself is faulty

- (a) arbitrary (unknown) value in val_{P_0}
 P_0 sends 1 to $\{P_1, P_3\}$, but 0 to P_2
 $\text{val}_{P_1} = \text{val}_{P_3} = 1$ but $\text{val}_{P_2} = 0$;
- (b) in P_1 $\text{om}(0, \{P_1, P_2, P_3\}, P_1)$
 in P_2 $\text{om}(0, \{P_1, P_2, P_3\}, P_2)$
 in P_3 $\text{om}(0, \{P_1, P_2, P_3\}, P_3)$
 forward correct in P_1, P_2 and P_3 ,
 but different values val_{P_i}



(c) majority

in P_1 mit $\langle 1, 0, 1 \rangle \implies 1$

$/^* \text{ values from } \langle P_1, P_2, P_3 \rangle ^*/$

in P_2 mit $\langle 1, 0, 1 \rangle \implies 1$

in P_3 mit $\langle 1, 0, 1 \rangle \implies 1$

$\implies \{P_1, P_2, P_3\}$ agree on the **same value 1**

Algorithms: $3m + 1$ Processes – m tolerable faults

► Lamport/Shostak/Pease 1982

$m + 1$ Exchange rounds; overall $\mathcal{O}(|PS|^m)$ Msgs

► Dolev/Reischuk: $2m + 1$ Rounds; $\mathcal{O}(|PS| * m + m^3)$ Msgs

► Gary/Moses (1993): $m + 1$ rounds; number of msgs polynomial

$\implies m$ tolerable faults require $m + 1$ deterministic exchanges

\implies Trade-Off: Number of messages vs. Number of Rounds

\implies *Algorithms are too costly for most situations*

ACM
ToPLaS
VI-105

JACM
1985

Fischer/
Lynch
1982

More Efficient Solution \implies Advanced System Model

► Use forgery-proof **Signatures** for all messages

► Protect message channels from eavesdropping and manipulation

\implies *Fakes when forwarding messages can be detected*

Expl.: Signatures allow even for 1-out-of-3 solution

Exchange $P_2 \longleftrightarrow P_3$ exposes inconsistent msgs from P_1

VI-104

End
VI.6.2

VI.6.3 Outlook: There are a lot more Algorithms

- Coordination in flat and nested transactions: 'Commit Protocols'
- Detecting and guaranteeing globally valid predicates
- Distributed Ledger Algorithms (BitCoin et al.), ...

Some Issues for further Studies

Complexity: Trade-off between preconditions, robustness and costs.

Correctness: Important issue for safety-critical distributed systems, but especially hard to tackle due to *State-Space explosion*.

Other System Models: Normally do not meet such favorable pre conditions, e.g.,

- * Indirect, even non-transitive, connectivity among nodes, e.g., in mobile or very heterogenous systems
- * Modern P2P systems: no stable structure and high churn rates

Conclusion – Distributed Algorithms

- ▶ **Overall Goal:** *Compensate typical deficits of distributed systems through additional algorithmic layers.*
- ◀ **To some extent achievable but there are limitations:**
 - ◁ Asynchronous systems with unpredictable error rates are barely manageable for practical applications.
 - ◀ Erroneous message channels are hard to overcome at all.
 - ◀ Constantly crashing nodes (processes) render productive work more or less impossible.
 - ⇒ both result in a '*trade-off*' between
 - **blocking** and even deadlocks due to long waiting periods
 - **life-locks** due to timeouts and permanent re-start of algorithms
 - ◁ Algorithms often costly, esp. number of msgs to be exchanged.

Practical Distributed System Development:

- * 'Hide' preconditions by using *Middleware* with QoS guarantees.
- * Confine algorithms to 'stable' settings, i.e., server environments.